

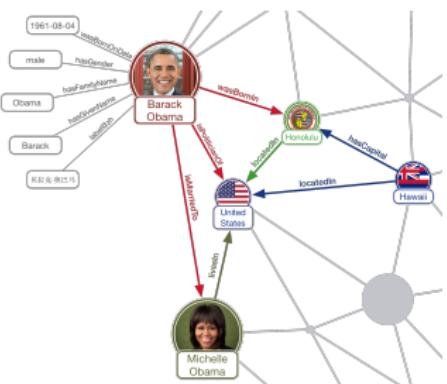
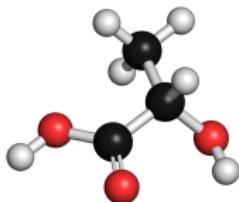
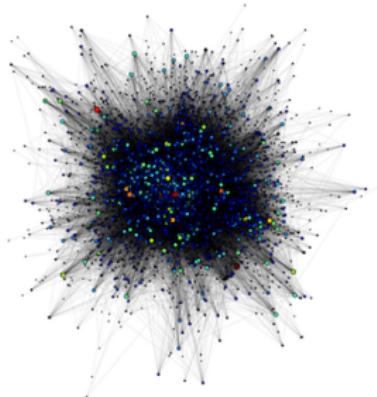
Machine Learning for Graphs based on Kernels @ DaSciM

M. Vazirgiannis & G. Nikolentzos

Data Science and Mining Team (DASCIM), LIX
École Polytechnique
Google Scholar: <https://bit.ly/2rwmvQU>
Twitter: @mvazirg

December, 2019

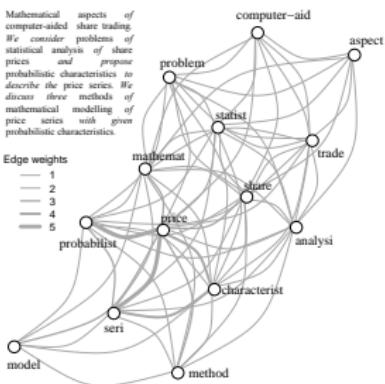
Graphs Are Everywhere



Mathematical aspects of computer-aided share trading. We consider problems of statistical analysis of share prices and propose probabilistic characteristics to describe the series. We discuss three methods of mathematical modelling of price series with given probabilistic characteristics.

Edge weights:

- 1
- 2
- 3
- 4
- 5



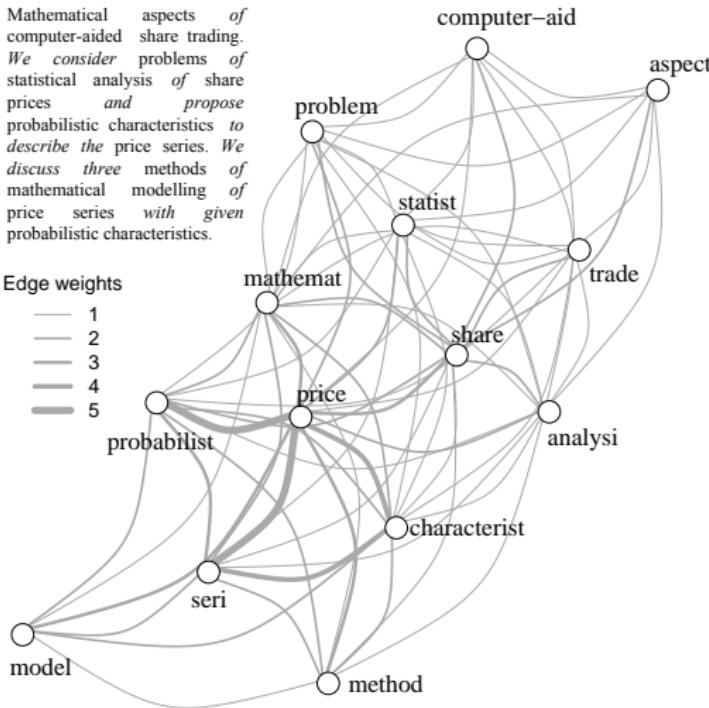
Why graphs?

Motivation - Text Categorization

Mathematical aspects of computer-aided share trading. We consider problems of statistical analysis of share prices and propose probabilistic characteristics to describe the price series. We discuss three methods of mathematical modelling of price series with given probabilistic characteristics.

Edge weights

- 1
- 2
- 3
- 4
- 5



Given a text, create a graph where

- vertices correspond to terms
- two terms are linked to each other if they co-occur within a fixed-size sliding window

Motivation - Text Categorization

Intuition: documents sharing same subgraphs belong to the same class

Given a set of documents and their graph representations:

Extract frequent subgraphs



- from the set of graphs

or

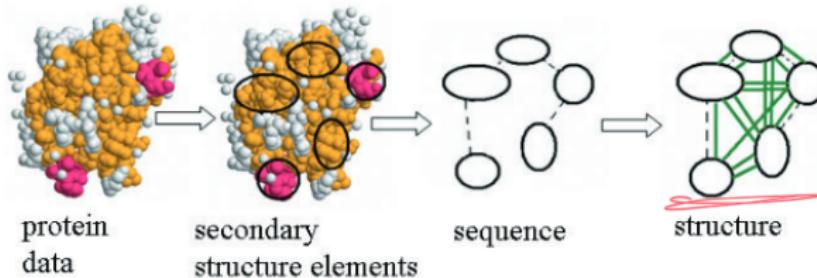
- from the set of the main cores of the graphs

Then, use frequent subgraphs as features for classification

Motivation - Protein Function Prediction

For each protein, create a graph that contains information about its

- structure
- sequence
- chemical properties



Use graph kernels to

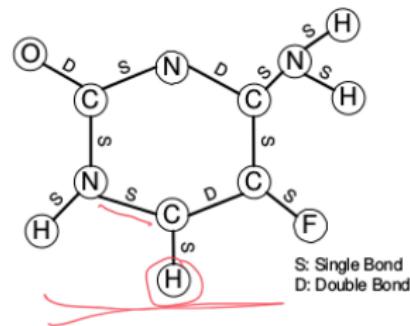
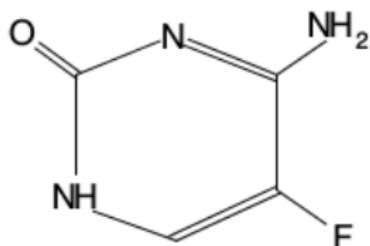


- measure structural similarity between proteins
- predict the function of proteins

Borgwardt et al. "Protein function prediction via graph kernels". Bioinformatics 21

Motivation - Chemical Compound Classification

Represent each chemical compound as a graph



Use a frequent subgraph discovery algorithm to discover the substructures that occur above a certain support constraint

Perform feature selection

Use the remaining substructures as features for classification

Deshpande et al. "Frequent substructure-based approaches for classifying chemical compounds". TKDE 17(8)

Motivation - Malware Detection

Given a computer program, create its control flow graph

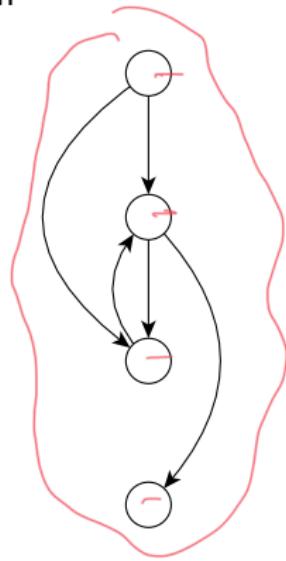
```
processed_pages.append(processed_page)
visited += 1
links = extract_links(html_code)
for link in links:
    if link not in visited_links:
        links_to_visit.append(link)

return create_vocabulary(processed_pages)

def parse_page(html_code):
    punct = re.compile("[^A-Za-z0-9]+")
    soup = BeautifulSoup(html_code, 'html.parser')
    text = soup.get_text()
    processed_text = punct.sub(" ", text)
    tokens = processed_text.split()
    tokens = [token.lower() for token in tokens]
    return tokens

def create_vocabulary(processed_pages):
    vocabulary = {}
    for processed_page in processed_pages:
        for token in processed_page:
            if token in vocabulary:
                vocabulary[token] += 1
            else:
                vocabulary[token] = 1

    return vocabulary
```



Compare the control flow graph of the problem against the set of control flow graphs of known malware

If it contains a subgraph isomorphic to these graphs → malicious code inside the program

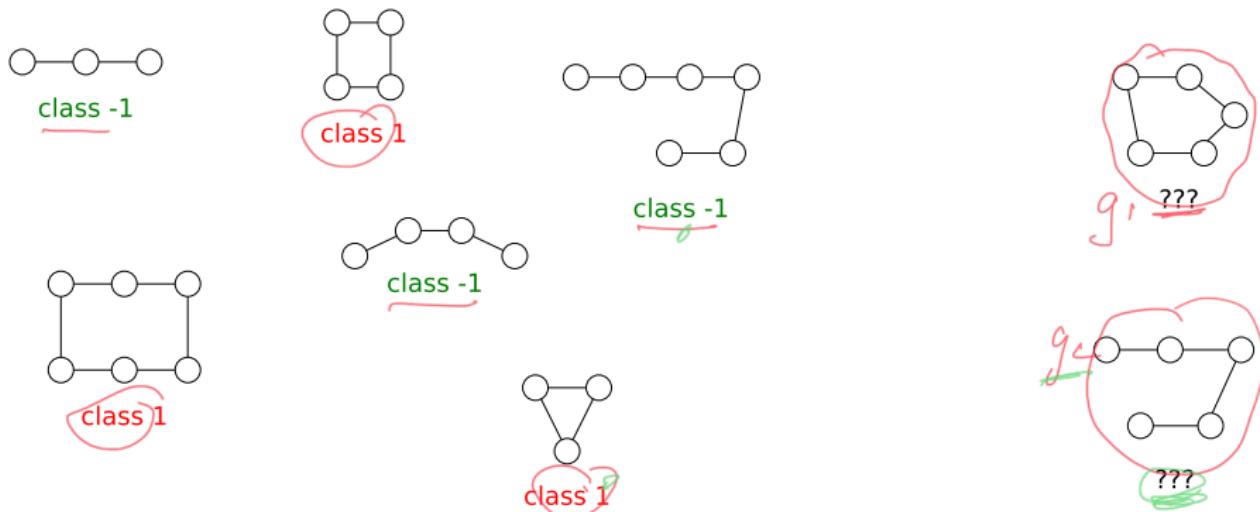
Gascon et al. "Structural detection of android malware using embedded call graphs". In AlSec'13

Machine Learning on Graphs

Machine learning tasks on graphs:

- Node classification: given a graph with labels on some nodes, provide a high quality labeling for the rest of the nodes 
- Graph clustering: given a graph, group its vertices into clusters taking into account its edge structure in such a way that there are many edges within each cluster and relatively few between the clusters
- Link Prediction: given a pair of vertices, predict if they should be linked with an edge 
- ✓ • Graph classification: given a set of graphs with known class labels for some of them, decide to which class the rest of the graphs belong 

Graph Classification



- Input data $G \in \mathcal{X}$
- Output $y \in \{-1, 1\}$
- Training set $\mathcal{D} = \{(G_1, y_1), \dots, (G_n, y_n)\}$
- Goal: estimate a function $f : \mathcal{X} \rightarrow \mathbb{R}$ to predict y from $f(x)$

Definition (Graph Comparison Problem)

Given two graphs G_1 and G_2 from the space of graphs \mathcal{G} , the problem of graph comparison is to find a mapping

$$s : \mathcal{G} \times \mathcal{G} \rightarrow \mathbb{R}$$

such that $s(G_1, G_2)$ quantifies the similarity of G_1 and G_2 .

Graph comparison is a topic of high significance

- It is the central problem for all learning tasks on graphs such as clustering and classification
- Most machine learning algorithms make decisions based on the similarities or distances between pairs of instances (e.g. k -nn)

Not an Easy Problem



Although graph comparison seems a tractable problem, it is very complex

[2015 Babai ($2^{O(\log(m))}$)
2019 Nikolentzos ($O(m^6)$)... (arxiv)]

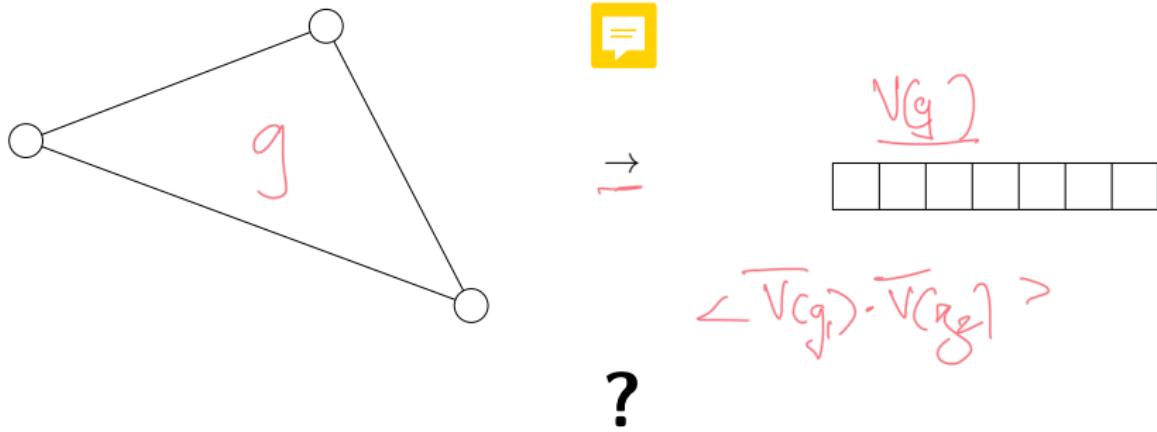
Many problems related to it are **NP-complete**

- subgraph isomorphism
- finding largest common subgraph

We are interested in algorithms capable of measuring the similarity between two graphs in **polynomial** time

Graphs to Vectors

- To analyze and extract knowledge from graphs, one needs to perform machine learning tasks
- Most machine learning algorithms require the input to be represented as a fixed-length feature vector
- There is no straightforward way to transform graphs to such a representation



What is a Kernel?

Definition (Kernel Function)

The function $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ is a kernel if it is:

- ① symmetric: $k(x, y) = k(y, x)$
- ② positive semi-definite: $\forall x_1, x_2, \dots, x_n \in \mathcal{X}$, the Gram Matrix \mathbf{K} defined by $\mathbf{K}_{ij} = k(x_i, x_j)$ is positive semi-definite

- If a function satisfies the above two conditions on a set \mathcal{X} , it is known that there exists a map $\phi : \mathcal{X} \rightarrow \mathcal{H}$ into a Hilbert space \mathcal{H} , such that:

$$k(x, y) = \langle \phi(x), \phi(y) \rangle$$

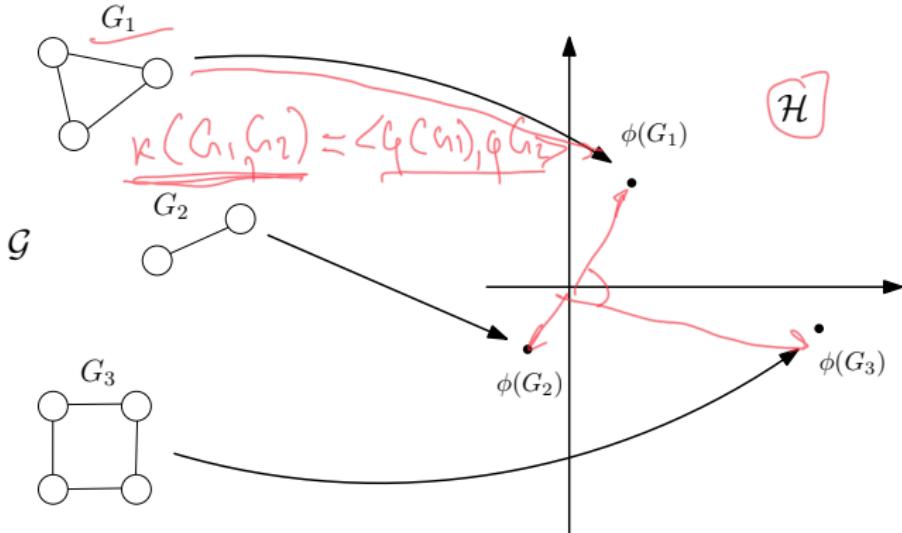
for all $(x, y) \in \mathcal{X}^2$ where $\langle \cdot, \cdot \rangle$ is the inner product in \mathcal{H}

- Informally, $k(x, y)$ is a measure of similarity between x and y

Definition (Graph Kernel)

A graph kernel $k : \mathcal{G} \times \mathcal{G} \rightarrow \mathbb{R}$ is a kernel function over a set of graphs \mathcal{G}

- It is equivalent to an inner product of the embeddings $\phi : \mathcal{X} \rightarrow \mathcal{H}$ of a pair of graphs into a Hilbert space
- Makes the whole family of kernel methods applicable to graphs



Kernel Trick

- Many machine learning algorithms can be expressed only in terms of inner products between vectors
- Let $\phi(G_1), \phi(G_2)$ be vector representations of graphs G_1, G_2 in a very high (possibly infinite) dimensional feature space
- Computing the explicit mappings $\phi(G_1), \phi(G_2)$ and their inner product $\langle \phi(x), \phi(y) \rangle$ for the pair of graphs can be computationally demanding
- The kernel trick avoids the explicit mapping by directly computing the inner product $\langle \phi(x), \phi(y) \rangle$ via the kernel function

Example

Let $\mathcal{X} = \mathbb{R}^2$ and

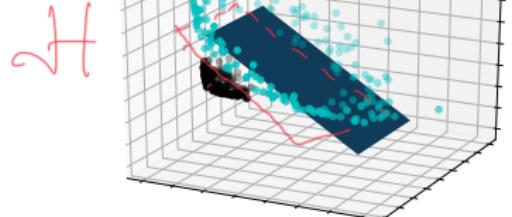
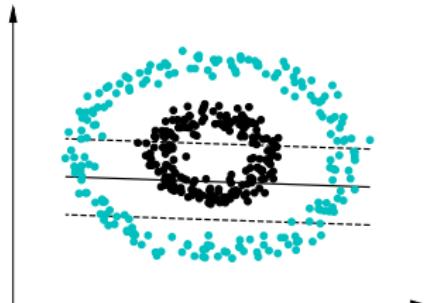
$$x = [x_1, x_2]^\top, y = [y_1, y_2]^\top \in \mathcal{X}$$

For any $x = [x_1, x_2]^\top$ let ϕ be a map
 $\phi : \mathbb{R}^2 \rightarrow \mathbb{R}^3$ defined as:

$$\phi(x) = [x_1^2, \sqrt{2}x_1x_2, x_2^2]^\top$$

Let also $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ a kernel defined
as $k(x, y) = \langle x, y \rangle^2$. Then

$$\begin{aligned} k(x, y) &= \langle x, y \rangle^2 \\ &= (x_1y_1 + x_2y_2)^2 \\ &= x_1^2y_1^2 + 2x_1y_1x_2y_2 + x_2^2y_2^2 \\ &= \langle \phi(x), \phi(y) \rangle \end{aligned}$$



Classification using SVM

- The standard SVM classifier addresses the following problem:
Given a set of N training objects along with their class labels
 $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N$, $x_i \in \mathcal{X}$, $y_i \in \mathcal{Y} = \{-1, +1\}$, learn a classifier $f : \mathcal{X} \rightarrow \mathcal{Y}$ that predicts the class labels of new objects
- SVM belongs to the family of large margin classifiers
↪ it seeks a hyperplane that separates the instances belonging to class -1 from those belonging to class 1
- This leads to the following dual optimization problem:

$$\begin{aligned} & \underset{\alpha}{\text{maximize}} \quad \sum_{i=1}^N \alpha_i - \frac{1}{4} \sum_{i=1}^N \sum_{j=1}^N \underbrace{\alpha_i \alpha_j y_i y_j}_{+} \langle \phi(x_i), \phi(x_j) \rangle \\ & \text{subject to} \quad \sum_{i=1}^N \underbrace{\alpha_i y_i}_{+} = 0 \\ & \qquad \qquad \qquad C \geq \alpha_i \geq 0 \quad \forall i \in \{1, \dots, N\} \end{aligned}$$

Graph Classification using SVM

- The standard SVM classifier addresses the following problem:
Given a set of N training objects along with their class labels
 $\mathcal{D} = \{(G_i, y_i)\}_{i=1}^N$, $G_i \in \mathcal{G}$, $y_i \in \mathcal{Y} = \{-1, +1\}$, learn a classifier $f : \mathcal{X} \rightarrow \mathcal{Y}$ that predicts the class labels of new objects
- SVM belongs to the family of large margin classifiers
→ it seeks a hyperplane that separates the instances belonging to class -1 from those belonging to class 1
- This leads to the following dual optimization problem:

$$\begin{aligned} & \underset{\alpha}{\text{maximize}} \quad \sum_{i=1}^N \alpha_i - \frac{1}{4} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y_i y_j \langle \phi(G_i), \phi(G_j) \rangle \\ & \text{subject to} \quad \sum_{i=1}^N \alpha_i y_i = 0 \\ & \quad C \geq \alpha_i \geq 0 \quad \forall i \in \{1, \dots, N\} \end{aligned}$$


Graph Classification using SVM

- The standard SVM classifier addresses the following problem:
Given a set of N training objects along with their class labels
 $\mathcal{D} = \{(G_i, y_i)\}_{i=1}^N$, $G_i \in \mathcal{G}$, $y_i \in \mathcal{Y} = \{-1, +1\}$, learn a classifier $f : \mathcal{X} \rightarrow \mathcal{Y}$ that predicts the class labels of new objects
- SVM belongs to the family of large margin classifiers
→ it seeks a hyperplane that separates the instances belonging to class -1 from those belonging to class 1
- This leads to the following dual optimization problem:

$$\begin{aligned} & \underset{\alpha}{\text{maximize}} \quad \sum_{i=1}^N \alpha_i - \frac{1}{4} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y_i y_j k(G_i, G_j) \\ & \text{subject to} \quad \sum_{i=1}^N \alpha_i y_i = 0 \\ & \quad C \geq \alpha_i \geq 0 \quad \forall i \in \{1, \dots, N\} \end{aligned}$$

Two Simple Kernels

The two kernels assume node/edge-labeled graphs

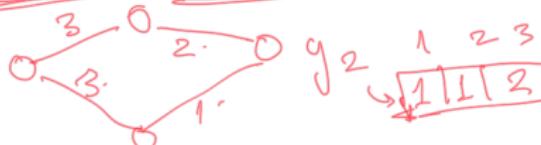
Vertex histogram kernel:

- The vertex label histogram of a graph G is a vector $f = [f_1, f_2, \dots, f_d]^\top$, such that $f_i = |\{v \in V : \ell(v) = i\}|$ for each $i \in \mathcal{L}$
- The vertex histogram kernel is then defined as:



$$\underline{k(G, G')} = \langle f, f' \rangle$$

Edge histogram kernel:

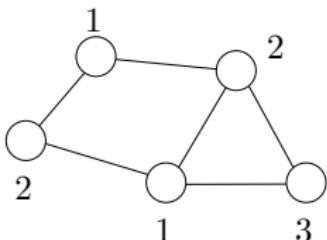


- The edge label histogram of a graph G is a vector $f = [f_1, f_2, \dots, f_d]^\top$, such that $f_i = |\{(v, u) \in E : \ell(v, u) = i\}|$ for each $i \in \mathcal{L}$.
- The edge histogram kernel is then defined as:

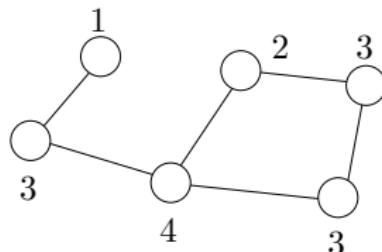
$$k(G, G') = \langle f, f' \rangle$$

Vertex Histogram Kernel

Example



G



G'

The vector representations of the two graphs are:

$$f_G = [2, 2, 1, 0]^\top$$



$$f_{G'} = [1, 1, 3, 1]^\top$$

Hence, the value of the kernel is:

$$k(G, G') = \langle f_G, f_{G'} \rangle = 7$$

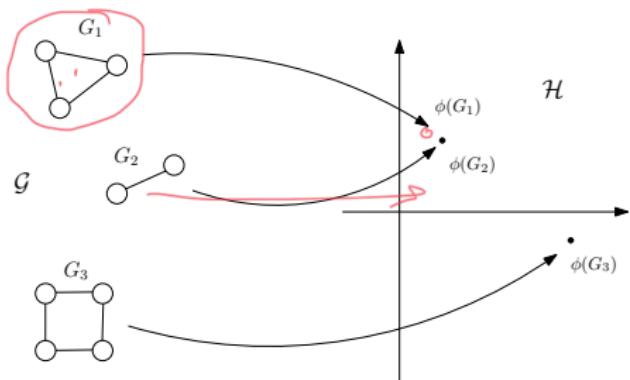
Expressiveness vs Efficiency

Complete Graph Kernels

Definition (Complete Graph Kernel)

A graph kernel $k(G_1, G_2) = \langle \phi(G_1), \phi(G_2) \rangle$ is complete if ϕ is injective

Hence, for complete graph kernels, $\phi(G_1) = \phi(G_2)$ iff G_1 and G_2 are isomorphic



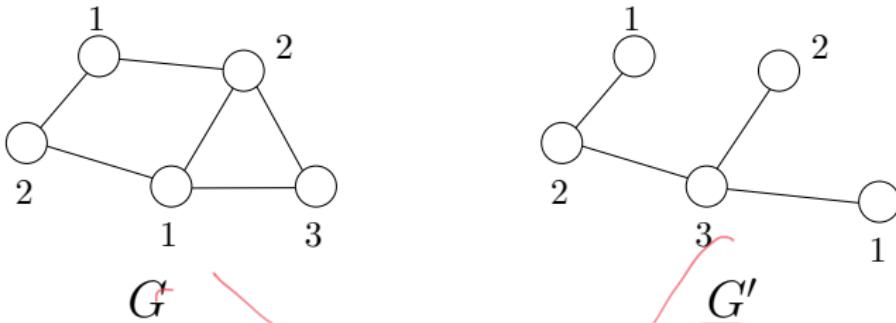
How hard is to compute a complete graph kernel?

Proposition

Computing any complete graph kernel is at least as hard as the graph isomorphism problem

Complete Graph Kernels

Clearly, the vertex and edge histogram kernels are not complete



The two graphs are not isomorphic. However

$$f_G = f_{G'} = [2, 2, 1, 0]^\top$$

$1 \approx 3A$



Expressiveness vs Efficiency

If the kernel is complete:

- Computation is at least as hard as the graph isomorphism problem
→ No polynomial algorithm for the graph isomorphism problem is known

If the kernel is not complete:

- It can be computed efficiently
- We can have $\phi(G_1) = \phi(G_2)$ even if $G_1 \not\cong G_2$
→ The kernel is not expressive enough



We are interested in kernels that can be computed in polynomial time (with small degree)

Expressive Power of Graph Kernels

Capitalize on concepts from property testing to measure the expressive power of graph kernels

Definition

A graph kernel *identifies* a property if no two graphs are mapped to the same feature vector unless they both have or both do not have the property (e.g., connected vs disconnected)

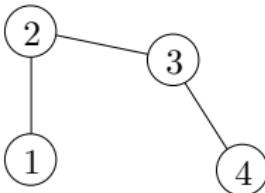
| Property \ Kernel | Weisfeiler-Lehman subtree kernel | Random Walk kernel | Shortest Path kernel | Graphlet kernel |
|-------------------|----------------------------------|--------------------|----------------------|-----------------|
| Connectivity | X | X | ✓ | X |
| Planarity | X | X | X | X |
| Bipartiteness | X | X | X | X |
| Triangle-freeness | X | X | X | ✓ |

Well-established kernels fail to identify fundamental properties

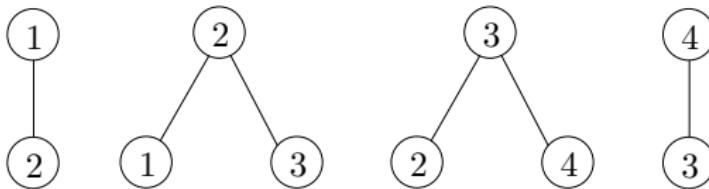
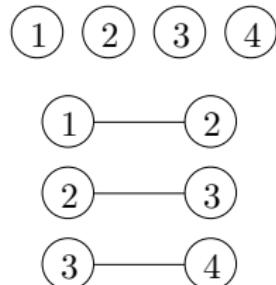
→ However, still they achieve state-of-the-art results on many datasets

Early Days of Graph Kernels

Convolution Kernels in a Nutshell



G



- Decompose structured objects into comparable parts
- Aggregate the values of similarity measures for individual parts

[Haussler. Tech Report'99]

- Let X be a set of composite objects (e.g., cars), and $\bar{X}_1, \dots, \bar{X}_D$ be sets of parts (e.g., wheels, brakes, etc.). All sets are assumed countable.
- Let R denote the relation “being part of”:

$$R(\bar{x}_1, \dots, \bar{x}_D, x) = 1, \text{ iff } \bar{x}_1, \dots, \bar{x}_D \text{ are parts of } x$$

- The inverse relation R^{-1} is defined as:

$$R^{-1}(x) = \{\bar{x} : R(\bar{x}, x) = 1\}$$

In other words, for each object x , $R^{-1}(x)$ is a set of component subsets, that are part of x

- We say that R is finite, if R^{-1} is finite for all $x \in X$

Example

| | |
|--|---|
| - x is a string | $x = \text{table}$ |
| - Subpart relation $R(\bar{x}_1, \bar{x}_2, x) = 1$ iff \bar{x}_1, \bar{x}_2 are (non-empty) strings such that $x = \text{concat}(\bar{x}_1, \bar{x}_2)$ | $\bar{x}_1 = t, \bar{x}_2 = \text{able}$ $\bar{x}_1 = \text{ta}, \bar{x}_2 = \text{ble}$ $\bar{x}_1 = \text{tab}, \bar{x}_2 = \text{le}$ $\bar{x}_1 = \text{tabl}, \bar{x}_2 = \text{e}$ |

Definition

Let $x, y \in X$ and \bar{x} and \bar{y} be the corresponding sets of parts. Let $K_d(\bar{x}_d, \bar{y}_d)$ be a kernel between the d -th parts of x and y ($1 \leq d \leq D$). Then the convolution kernel between x and y is defined as:

$$K(x, y) = \sum_{\bar{x} \in R^{-1}(x)} \sum_{\bar{y} \in R^{-1}(y)} \prod_{d=1}^D K_d(x_d, y_d)$$

Substructures-based Kernels

A large number of graph kernels compare substructures of graphs that are computable in polynomial time:

- walks
 - shortest paths
 - subtree patterns
 - graphlets
- ⋮

These kernels are instances of the R-convolution framework

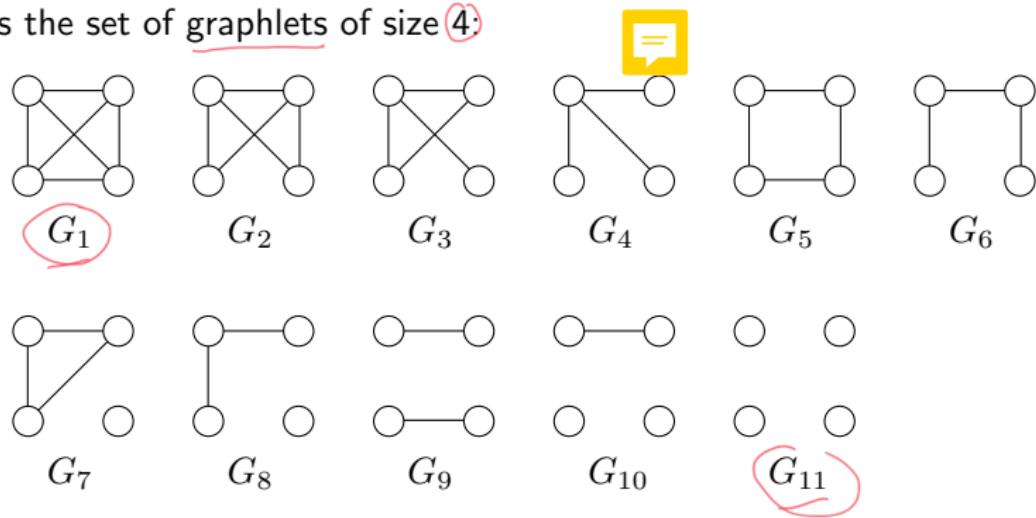
Graphlet Kernel

The graphlet kernel compares graphs by counting *graphlets*

A graphlet corresponds to a small subgraph

- typically of 3,4 or 5 vertices

Below is the set of graphlets of size 4:



[Shervashidze et al., AISTATS'09]

Graphlet Kernel

Let $\mathcal{G} = \{\text{graphlet}_1, \text{graphlet}_2, \dots, \text{graphlet}_r\}$ be the set of size- k graphlets 

Let also $f_G \in \mathbb{N}^r$ be a vector such that its i -th entry is $f_{G,i} = \#(\text{graphlet}_i \sqsubseteq G)$

The graphlet kernel is defined as:

$$k(G_1, G_2) = \langle f_{G_1}, f_{G_2} \rangle$$

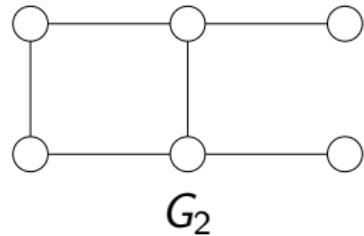
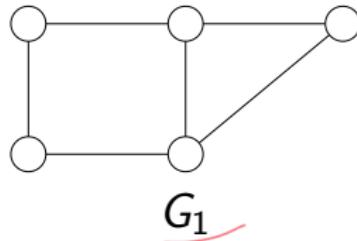
Problems:

- There are $\binom{n}{k}$ size- k subgraphs in a graph
- Exhaustive enumeration of graphlets is very expensive

Requires $O(n^k)$ time 

- For labeled graphs, the number of graphlets increases further

Example



The vector representations of the graphs above according to the set of graphlets of size 4 is:

$$\begin{aligned} \text{--- } f_{G_1} &= [0, 0, 2, 0, 1, 2, 0, 0, 0, 0, 0]^\top \quad \text{💡} \\ \text{--- } f_{G_2} &= [0, 0, 0, 2, 1, 5, 0, 4, 0, 3, 0]^\top \quad \text{💡} \end{aligned}$$

Hence, the value of the kernel is:

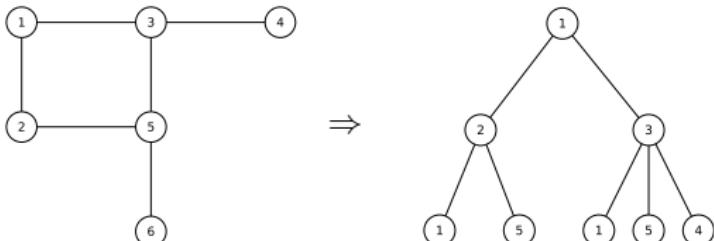
$$k(G_1, G_2) = \langle f_{G_1}, f_{G_2} \rangle = 11$$

Subtree Kernel

Compares subtree patterns in two graphs

A subtree pattern is a subgraph of a graph which has

- a root vertex
- no cycles



Subtree of height 2 rooted at vertex 1

The height of a subtree is the maximum distance between the root and any other node in the subtree

If there are cycles in the graph, a vertex can appear more than once in a subtree pattern

- it is treated as a distinct vertex such that the pattern is still a cycle-free tree

For all pairs of nodes v from G_1 and u from G_2 :

- Create the subtree patterns of height h rooted at v, u
- Compare v and u via a kernel function
- Recursively compare all vertices of the subtree patterns of v and u via a kernel function

[Ramon and Gartner. MGTS'03]

Subtree Kernel

Given a pair of graphs $G_1 = (V_1, E_1)$ and $G_2 = (V_2, E_2)$, the subtree kernel of height h is defined as:

$$k(G_1, G_2) = \sum_{v_1 \in V_1} \sum_{v_2 \in V_2} k_h(v_1, v_2)$$

where

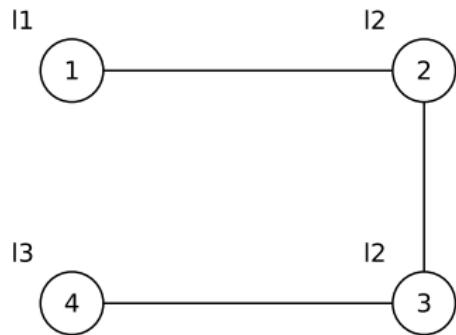
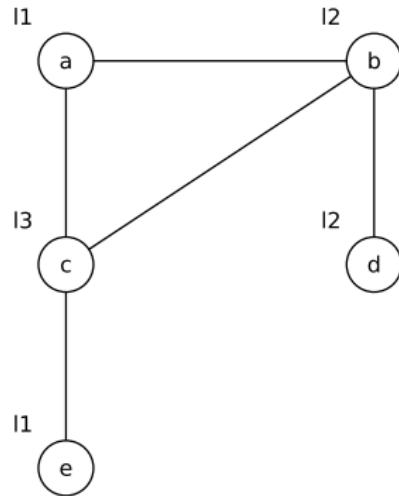
$$k_h(v_1, v_2) = \begin{cases} \delta(\ell(v_1) = \ell(v_2)) & \text{if } h = 0 \\ \lambda_{v_1} \lambda_{v_2} \delta(\ell(v_1) = \ell(v_2)) \sum_{R \in \mathcal{M}(v_1, v_2)} \prod_{(w_1, w_2) \in R} k_{h-1}(w_1, w_2) & \text{if } h > 0 \end{cases}$$

where $\delta(\cdot, \cdot)$ is the Kronecker delta function that equals 1 if its arguments are equal, 0 otherwise, λ_{v_1} and λ_{v_2} are weights associated with nodes v_1 and v_2 , and

$$\begin{aligned} \mathcal{M}(v_1, v_2) = & \left\{ R \subseteq \mathcal{N}(v_1) \times \mathcal{N}(v_2) \mid (\forall (u_1, u_2), (w_1, w_2) \in R : u_1 = w_1 \Leftrightarrow u_2 = w_2) \right. \\ & \left. \wedge (\forall (u_1, u_2) \in R : \ell(u_1) = \ell(u_2)) \right\} \end{aligned}$$

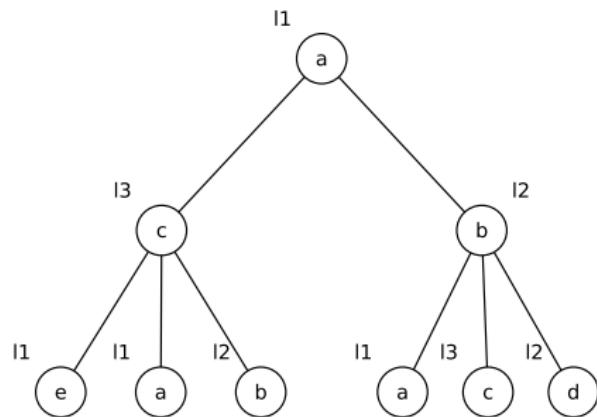
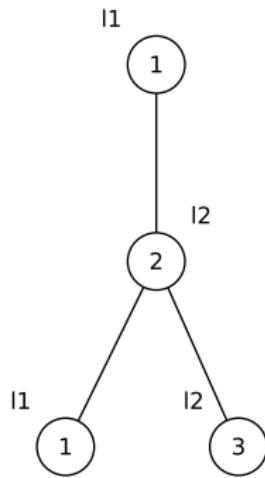
Example

We are given the following graphs


$$G_1$$

$$G_2$$

Example

Below are given the subtrees of G_1 and G_2 with height 2 rooted at 1 and a respectively



We will compute $k_2(1, a)$

Example

We set $\lambda_v = 1$ for all $v \in V_1 \cup V_2$ and we have:

$$k_2(1, a) = \delta(\ell(1) = \ell(a)) \sum_{R \in \mathcal{M}(1, a)} \prod_{(v_1, v_2) \in R} k_1(v_1, v_2)$$

- $\delta(\ell(1) = \ell(a)) = 1$ since $\ell(1) = \ell(a) = l1$
- $\mathcal{M}(1, a) = \{(2, b)\}$ since $\ell(2) = \ell(b) = l2$

Hence, we will next compute $k_1(2, b)$

$$k_1(2, b) = \delta(\ell(2) = \ell(b)) \sum_{R \in \mathcal{M}(2, b)} \prod_{(v_1, v_2) \in R} k_0(v_1, v_2)$$

- $\delta(\ell(2) = \ell(b)) = 1$ since $\ell(2) = \ell(b) = l2$
- $\mathcal{M}(2, b) = \{(1, a), (3, d)\}$ since $\ell(1) = \ell(a) = l1$ and $\ell(3) = \ell(d) = l2$

Example

At height 0, we have:

$$k_0(1, a) = k_0(3, d) = 1$$

Therefore,

$$k_1(2, b) = k_0(1, a)k_0(3, d) = 1$$

And finally,

$$k_2(1, a) = k_1(2, b) = 1$$

Subtree kernel

Pros: Richer representation of graph structure



Cons: Very high complexity

- $\mathcal{O}(n^2 h 4^d)$ where d is the maximum degree of the pair of graphs

Shortest Path Kernel

Compares the length of shortest-paths of two graphs

- and their endpoints in labeled graphs

Floyd-transformation

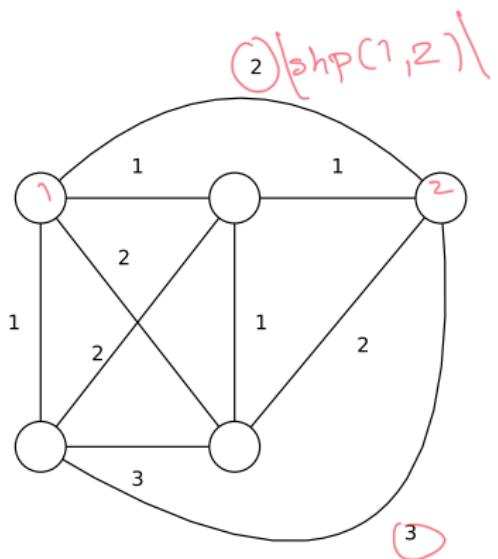
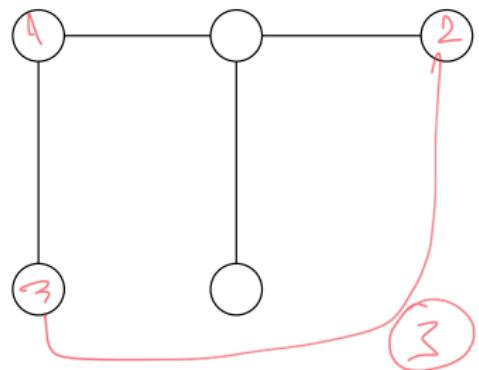
Transforms the original graphs into shortest-paths graphs

- Compute the shortest-paths between all pairs of vertices of the input graph G using some algorithm (i. e. Floyd-Warshall)
- Create a shortest-path graph S which contains the same set of nodes as the input graph G
- All nodes which are connected by a walk in G are linked with an edge in S
- Each edge in S is labeled by the shortest distance between its endpoints in G

[Borgwardt and Kriegel. ICDM'05]

Example

Floyd-transformation



G

S



Shortest Path Kernel

Given the Floyd-transformed graphs $S_1 = (V_1, E_1)$ and $S_2 = (V_2, E_2)$ of G_1 and G_2 , the shortest path kernel is defined as:

$$k(G_1, G_2) = \sum_{e_1 \in E_1} \sum_{e_2 \in E_2} k_{\text{edge}}(e_1, e_2)$$

where k_{edge} is a kernel on edges

- For unlabeled graphs, it can be:



$$k_{\text{edge}}(e_1, e_2) = \delta(\ell(e_1), \ell(e_2)) = \begin{cases} 1 & \text{if } \ell(e_1) = \ell(e_2), \\ 0 & \text{otherwise} \end{cases}$$

where $\ell(e)$ gives the label of edge e

- For labeled graphs, it can be:

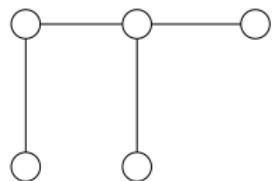


$$k_{\text{edge}}(e_1, e_2) = \begin{cases} 1 & \text{if } \ell(e_1) = \ell(e_2) \wedge \ell(e_1^1) = \ell(e_2^1) \wedge \ell(e_1^2) = \ell(e_2^2), \\ 0 & \text{otherwise} \end{cases}$$

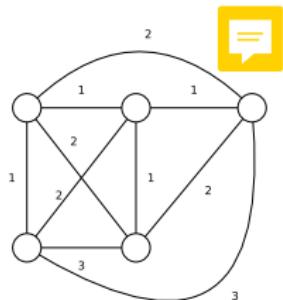
where e^1, e^2 are the two endpoints of e

Example

Floyd-transformations

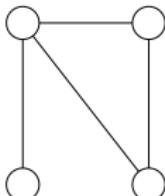


⇒

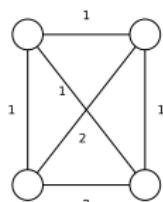


G_1

S_1



⇒



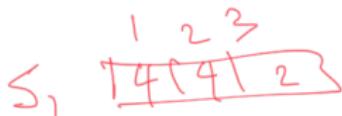
G_2

S_2

Example

In S_1 we have:

- 4 edges with label 1
- 4 edges with label 2
- 2 edges with label 3



$$K(G_1, G_2) = \overline{S}_1 \cdot \overline{S}_2 = 24$$

In S_2 we have:

- 4 edges with label 1
- 2 edges with label 2



Hence, the value of the kernel is:

$$k(G_1, G_2) = \sum_{e_1 \in E_1} \sum_{e_2 \in E_2} k_{\text{edge}}(e_1, e_2) = 4 \cdot 4 + 4 \cdot 2 = 24$$

Computing the shortest path kernel includes:

- Computing shortest paths for all pairs of vertices in the two graphs: $\mathcal{O}(n^3)$
- Comparing all pairs of shortest paths from the two graphs: $\mathcal{O}(n^4)$

Hence, runtime is $\mathcal{O}(n^4)$

Problems:

- Very high complexity for large graphs
- Shortest-path graphs may lead to memory problems on large graphs

Cyclic Pattern Kernel

The cyclic pattern kernel

- decomposes a graph into cyclic and tree patterns
- counts the number of common patterns which occur in two graphs

Cycles:

- Let $\mathcal{S}(G)$ denote the set of cycles of a graph G
- Let also $\pi(C)$ denote the canonical representation of a cycle C
- The set of cyclic patterns of G is defined by $\mathcal{C}(G) = \{\pi(C) : C \in \mathcal{S}(G)\}$

Trees:

- By removing all the edges of all cycles, the kernel obtains a set of trees
- The kernel computes the canonical representation $\pi(T)$ of each tree T
- The set of tree patterns of G is then defined by $\mathcal{T}(G) = \{\pi(T) : T \text{ is a tree}\}$

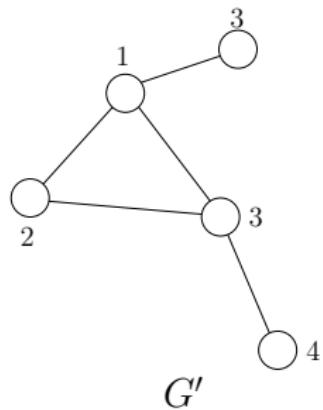
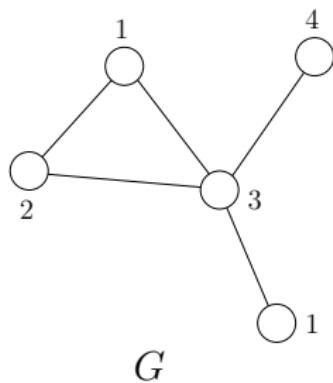
The cyclic pattern kernel is then defined as

$$k(G, G') = |\mathcal{C}(G) \cap \mathcal{C}(G')| + |\mathcal{T}(G) \cap \mathcal{T}(G')|$$

Problems:

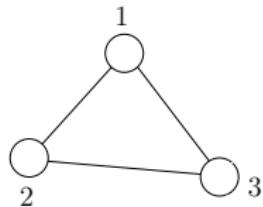
- Number of cyclic and tree patterns can be exponential in the number of vertices n
- Computing the cyclic pattern kernel on general graphs is NP-hard
- Can only be applied to graphs where the number of cycles is polynomially bounded
[Horvath et al., KDD'04]

Example

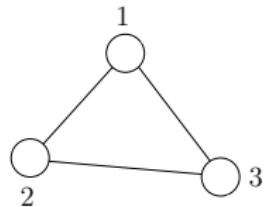


Extract cyclic and tree patterns from G, G'

Example



C_1



C'_1

$$\mathcal{C}(G) = \{\pi(C_1)\} = \{(1, 2, 3)\}$$

$$\mathcal{C}(G') = \{\pi(C'_1)\} = \{(1, 2, 3)\}$$

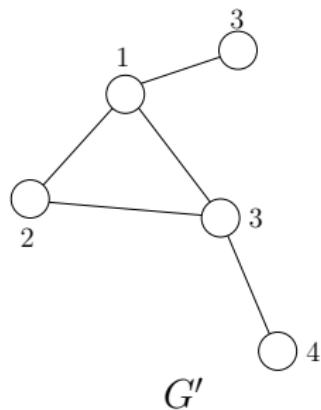
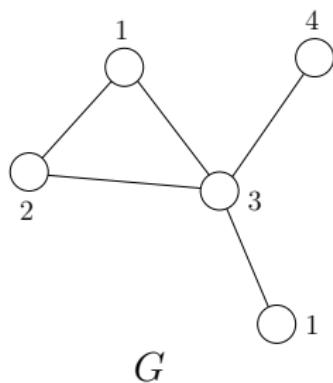
Example



$$\mathcal{T}(G) = \{\pi(T_1)\} = \{(1, 3, 4)\}$$

$$\mathcal{T}(G') = \{\pi(T'_1), \pi(T'_2)\} = \{(1, 3), (3, 4)\}$$

Example



Hence, kernel equal to

$$k(G, G') = |\mathcal{C}(G) \cap \mathcal{C}(G')| + |\mathcal{T}(G) \cap \mathcal{T}(G')| = 1$$

Random Walk Kernel

- Probably the most well-studied family of graph kernels
- Counts matching walks in two graphs

Product graph

Given two graphs $G_1 = (V_1, E_1)$ and $G_2 = (V_2, E_2)$, their direct product G_\times is a graph with vertex set:

$$V_\times = \{(v_1, v_2) : v_1 \in V_1, v_2 \in V_2\} \text{ for unlabeled graphs}$$

or

$$V_\times = \{(v_1, v_2) : v_1 \in V_1, v_2 \in V_2, \ell(v1) = \ell(v2)\} \text{ for labeled graphs}$$

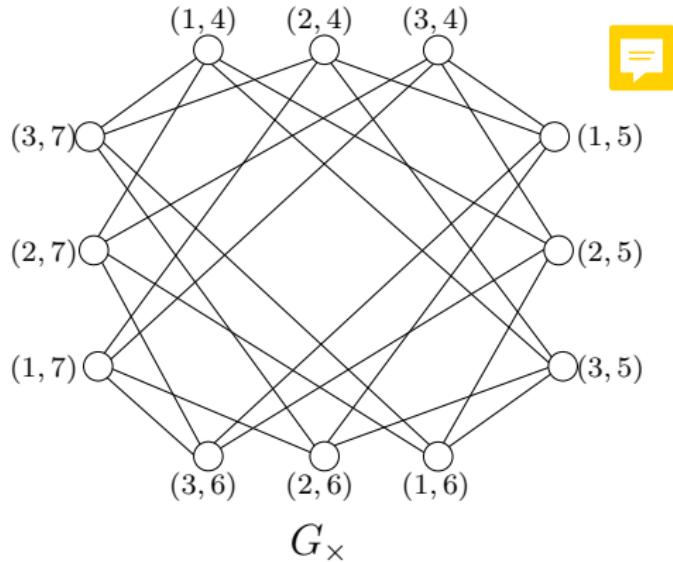
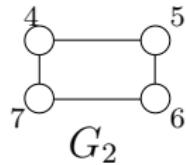
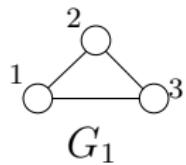
and edge set:

$$E_\times = \{((v_1, v_2), (u_1, u_2)) : (v_1, u_1) \in E_1, (v_2, u_2) \in E_2\}$$

- vertices: pairs of vertices from G_1 and G_2
- draw edge if corresponding vertices of G_1 and G_2 are adjacent in G_1 and G_2

[Gartner et al., COLT/Kernel'03]

Example



Random Walk Kernel

The k -th power of the adjacency matrix A of G computes walks of length k
 $\hookrightarrow A_{ij}^k = \text{number of walks of length } k \text{ from vertex } i \text{ to vertex } j$

Performing a random walk on G_x is equivalent to performing a simultaneous random walk on G_1 and G_2 of length k

- Common walks of length k can be computed using A_x^k

For $k \in \mathbb{N}$, the k -step random walk kernel is defined as:

$$K_x^k(G_1, G_2) = \sum_{i,j=1}^{|V_x|} \left[\sum_{r=0}^k \lambda_r A_x^r \right]_{ij}$$

where $\lambda_0, \lambda_1, \dots, \lambda_k$ positive weights and $A_x^0 = I$

Random Walk Kernel

For $k \rightarrow \infty$, we obtain the geometric random walk kernel $K_{\times}^{\infty}(G_1, G_2)$

If $\lambda_r = \lambda^r$, $K_{\times}^{\infty}(G_1, G_2)$ can be directly computed as follows:

$$K_{\times}^{\infty}(G_1, G_2) = \sum_{i,j=1}^{|V_{\times}|} \left[\sum_{r=0}^{\infty} \lambda^r A_{\times}^r \right]_{ij} = e^{\top} (I - \lambda A_{\times})^{-1} e$$

where e the all-ones vector

Problem: computational complexity is $\mathcal{O}(n^6)$

Solution: Efficient computation (almost $\mathcal{O}(n^3)$) using:

- Sylvester equations
- Conjugate gradient solver
- Fixed-point iterations
- Spectral decompositions
- λ should be non greater than the largest eigenvalue of A_{\times}

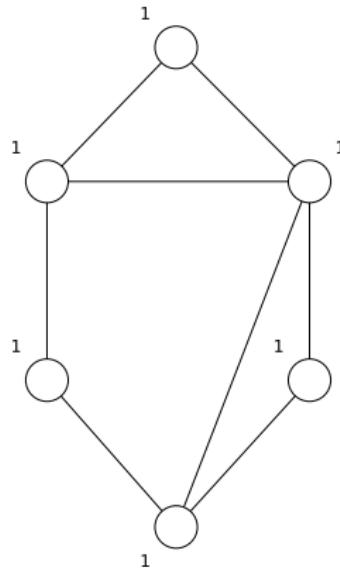
[Vishwanathan et al., JMLR 11.Apr (2010)]

Neighborhood Aggregation Approaches

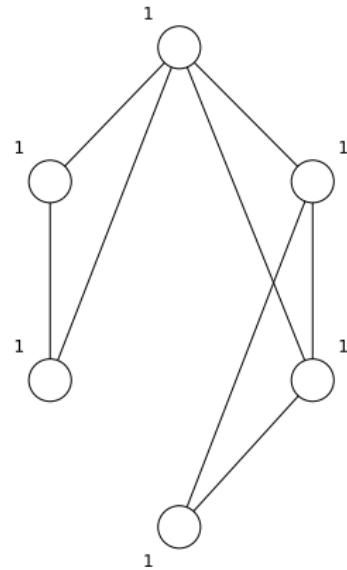
Weisfeiler-Lehman Test of Isomorphism

May answer if two graphs are not isomorphic

Run the Weisfeiler-Lehman algorithm for the following pair of graphs



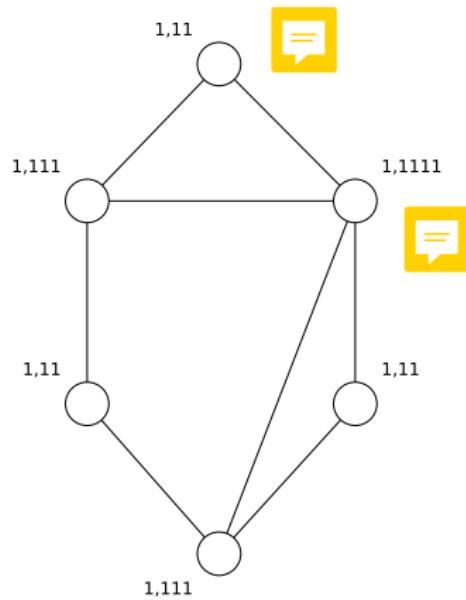
G_1



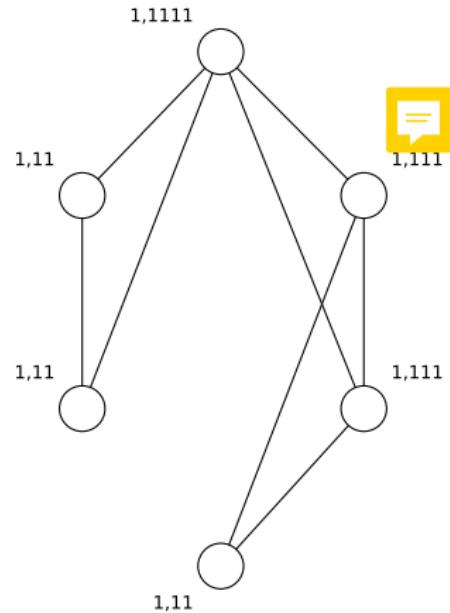
G_2

Iteration 1

First step: Augment the labels of the vertices by the sorted set of labels of neighbouring vertices



G_1



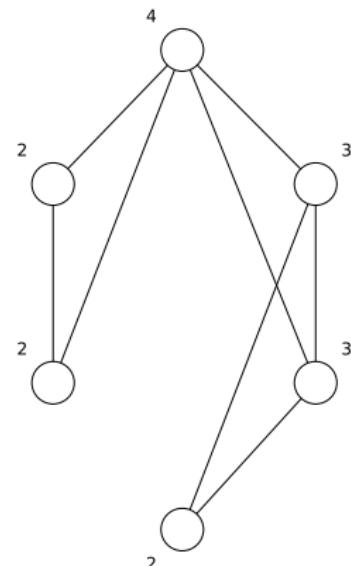
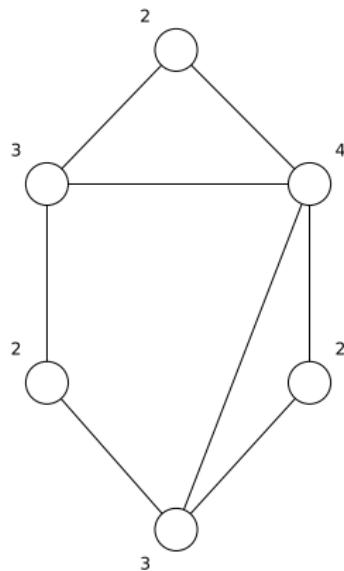
G_2

Iteration 1

Second step: Compress the augmented labels into new, short labels:

- o $1, 11 \rightarrow 2$
- o $1, 111 \rightarrow 3$

- o $1, 1111 \rightarrow 4$

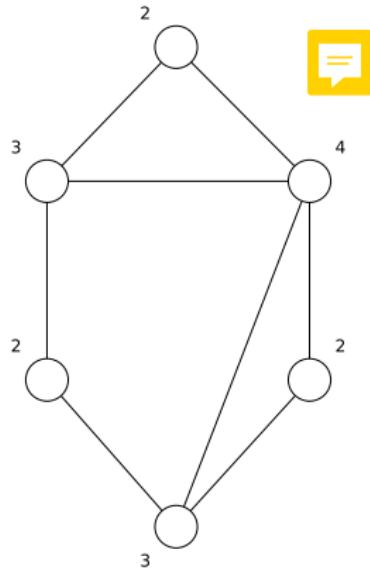


G_1

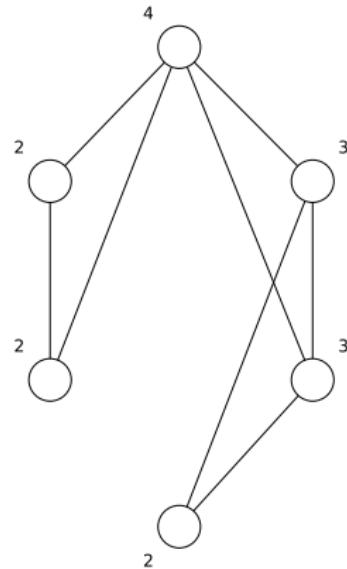
G_2

Iteration 1

Are the label sets of G_1 and G_2 identical?



G_1



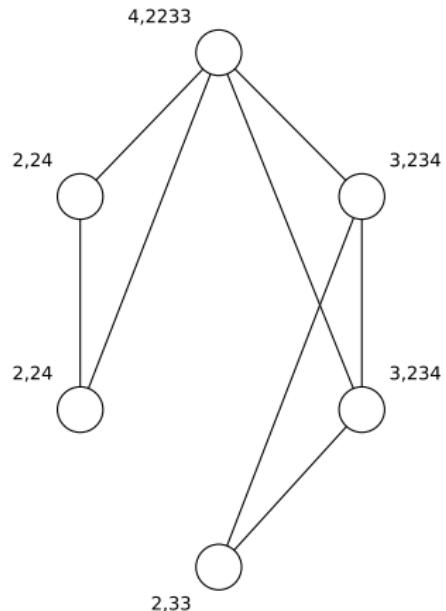
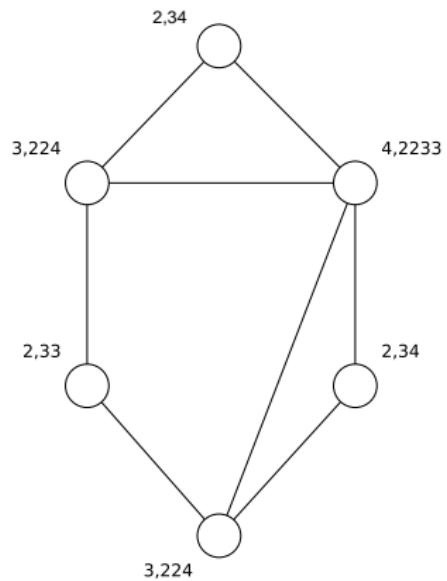
G_2

Yes!!!

Continue to the next iteration

Iteration 2

First step: Augment the labels of the vertices by the sorted set of labels of neighbouring vertices



G_1

G_2

Iteration 2

Second step: Compress the augmented labels into new, short labels:

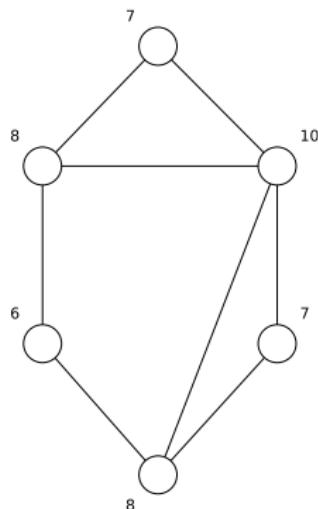
- o $2, 24 \rightarrow 5$

- o $2, 33 \rightarrow 6$

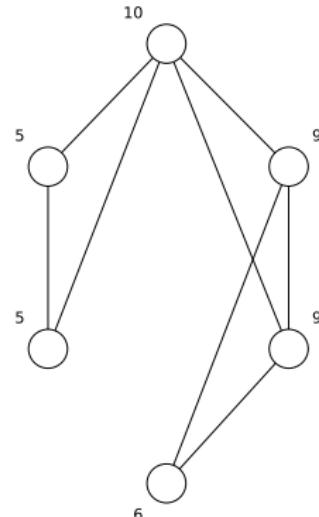
- o $2, 34 \rightarrow 7$

- o $3, 234 \rightarrow 9$

- o $4, 2233 \rightarrow 10$



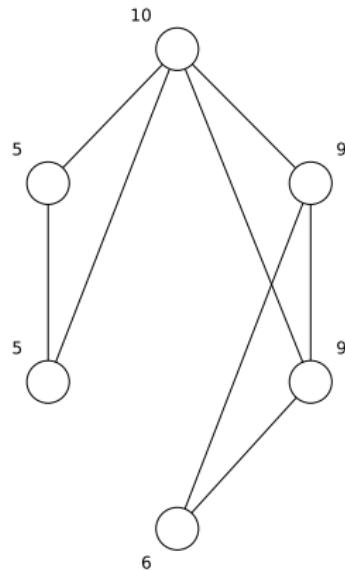
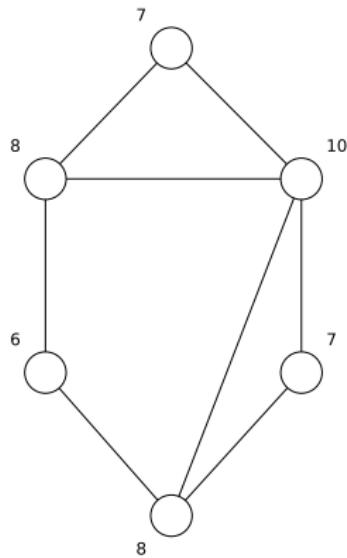
G_1



G_2

Iteration 2

Are the label sets of G_1 and G_2 identical?



G_1

G_2

No!!!

Graphs are not isomorphic

Weisfeiler-Lehman Framework

Let G^1, G^2, \dots, G^h be the graphs emerging from graph G at the iteration $1, 2, \dots, h$ of the Weisfeiler-Lehman algorithm

Then, the Weisfeiler-Lehman kernel is defined as:



$$k_{WL}^h(G_1, G_2) = k(G_1, G_2) + k(G_1^1, G_2^1) + k(G_1^2, G_2^2) + \dots + k(G_1^h, G_2^h)$$

where $k(\cdot, \cdot)$ is a base kernel (e.g. subtree kernel, shortest path kernel, ...)

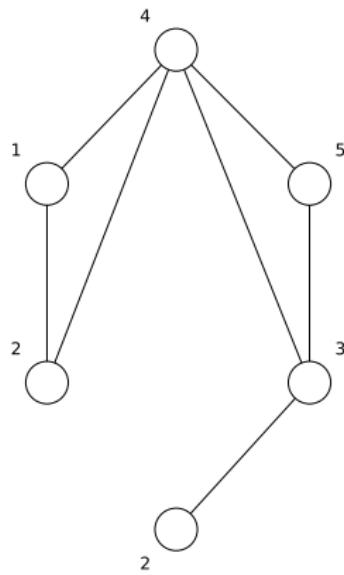
At each iteration of the Weisfeiler-Lehman algorithm:

- run a graph kernel for labeled graphs
- the new kernel values are added to the ones of the previous iteration

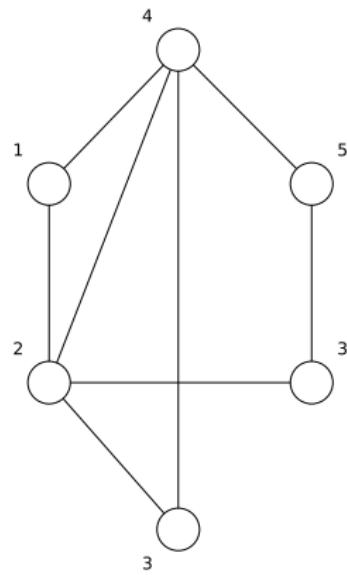
[Shervashidze et al., JMLR 12.Sep (2011)]

Weisfeiler-Lehman Subtree Kernel

Counts matching pairs of labels in two graphs after each iteration



G_1

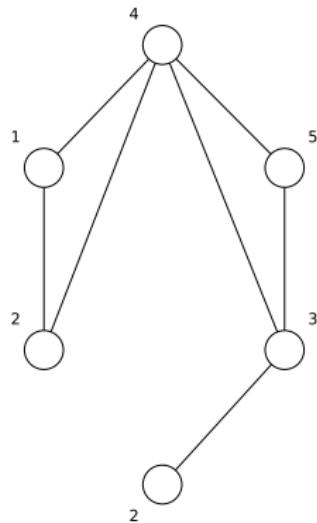


G_2

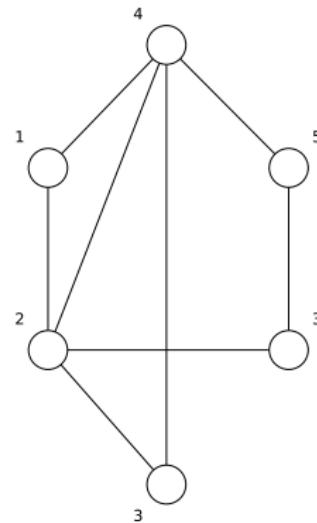
Initialization

Feature vector for a graph G :

$$\phi(G) = \{\#\text{nodes with label 1}, \#\text{nodes with label 2}, \dots, \#\text{nodes with label } l\}$$



G_1

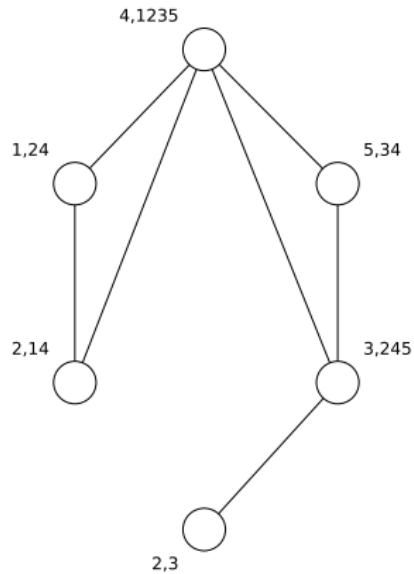


G_2

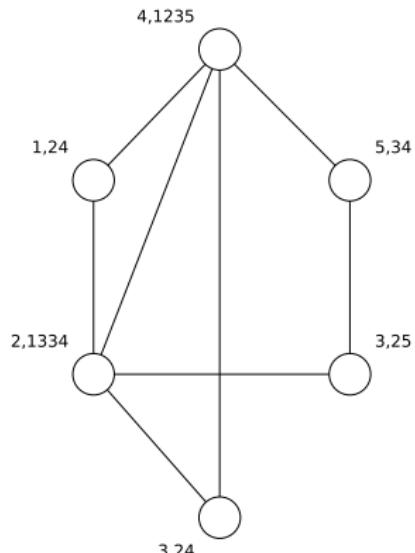
$$\begin{aligned}\phi(G_1) &= [1, 2, 1, 1, 1]^\top & \phi(G_2) &= [1, 1, 2, 1, 1]^\top \\ k(G_1, G_2) &= \langle \phi(G_1), \phi(G_2) \rangle = 7\end{aligned}$$

Iteration 1

First step: Augment the labels of the vertices by the sorted set of labels of neighbouring vertices



G_1

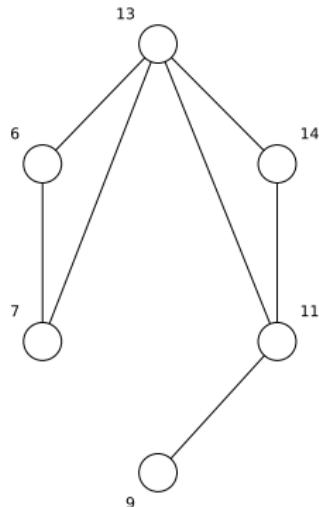


G_2

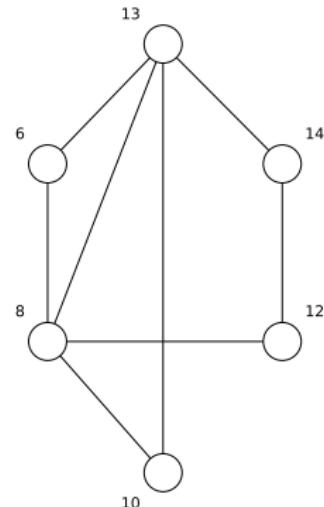
Iteration 1

Second step: Compress the augmented labels into new, short labels:

- o $1, 24 \rightarrow 6$
- o $2, 14 \rightarrow 7$
- o $2, 1334 \rightarrow 8$
- o $2, 3, 24 \rightarrow 9$
- o $3, 24 \rightarrow 10$
- o $3, 245 \rightarrow 11$
- o $3, 25 \rightarrow 12$
- o $4, 1235 \rightarrow 13$
- o $5, 34 \rightarrow 14$



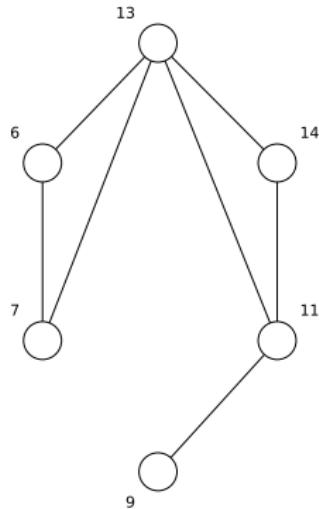
G_1



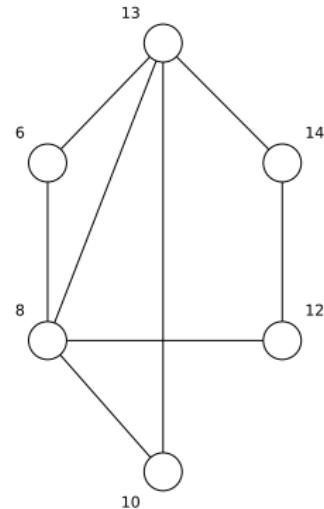
G_2

Iteration 1

Third step: Compute kernel value for iteration $h = 1$ and add it to previous kernel value



G_1



G_2

$$\phi(G_1^1) = [1, 1, 0, 1, 0, 1, 0, 1, 1]^\top \quad \phi(G_2^1) = [1, 0, 1, 0, 1, 0, 1, 1, 1]^\top$$

$$k(G_1^1, G_2^1) = \langle \phi(G_1^1), \phi(G_2^1) \rangle = 3$$

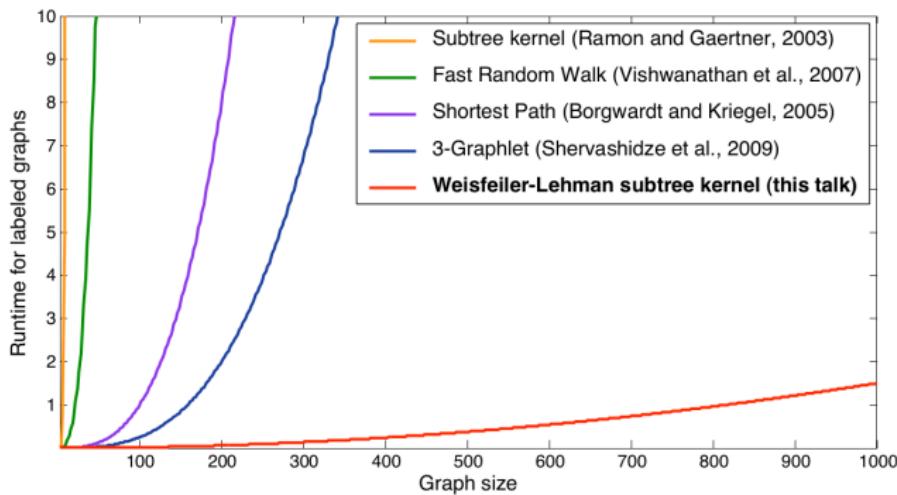
$$k_{WL}^1(G_1, G_2) = k(G_1, G_2) + k(G_1^1, G_2^1) = 10$$

Weisfeiler-Lehman Subtree Kernel

Computing the Weisfeiler-Lehman Subtree Kernel takes $\mathcal{O}(hm)$ time

- very efficient

Comparison to other well-known kernels



More Recent Approaches

Lovász ϑ kernel

Compares graphs based on the orthonormal representation associated with the Lovász number

- the orthonormal representation captures **global** graph properties

Orthonormal representation of a graph $G = (V, E)$:

- each vertex $i \in V$ is assigned a unit vector u_i , $\|u_i\| = 1$
- let $U_G = \{u_1, u_2, \dots, u_n\}$ be the set of all vectors
- for $i, j \in V$, if $(i, j) \notin E$, then $u_i^\top u_j = 0$



An interesting orthonormal representation is associated with the Lovász number $\vartheta(G)$

Definition (Lovász number)

For a graph $G = (V, E)$,

$$\vartheta(G) = \min_{c, U_G} \max_{i \in V} \frac{1}{(c^\top u_i)^2}$$

where the minimization is taken over all orthonormal representations U_G and all unit vectors c

[Johansson et al., ICML'14]

Lovász ϑ kernel

Given a subset of vertices $S \subseteq V$, the Lovász value of the subgraph induced by S is:

$$\vartheta_S(G) = \min_c \max_{u_i \in U_{G|S}} \frac{1}{(c^\top u_i)^2}$$

where $U_{G|S} = \{u_i \in U_G : i \in S\}$

The Lovász kernel is then defined as:

$$k_\vartheta(G_1, G_2) = \sum_{\substack{S_1 \subseteq V_1 \\ |S_1|=|S_2|}} \sum_{S_2 \subseteq V_2} \frac{1}{Z} k(\vartheta_{S_1}(G_1), \vartheta_{S_2}(G_2))$$

where $Z = \binom{n_1}{d} \binom{n_2}{d}$, $d = |S_1| = |S_2|$ and $k(\cdot, \cdot)$ is a base kernel (e.g. linear, gaussian)

Problem: Computing the Lovász ϑ kernel is very expensive since → requires computing the Lovász value for all subgraphs of the two graphs

Solution: Sampling

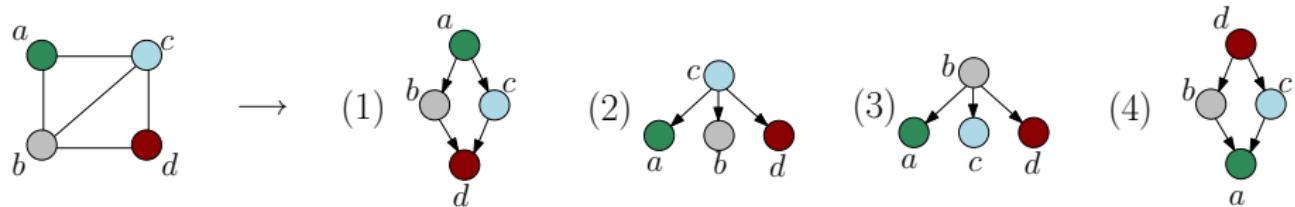
→ Evaluate the Lovász value for a smaller number of subgraphs of size d

Ordered Decomposition DAGs Kernel

General idea:

- Decomposes graphs into multisets of directed acyclic graphs (DAGs)
- Uses existing tree kernels to compare these DAGs

Generates one unordered rooted DAG for each vertex (keeps only edges belonging to the shortest paths)



Then, the kernel is defined as:

$$k(G, G') = \sum_{D \in DD(G)} \sum_{D' \in DD(G')} k_{DAG}(D, D')$$

where $DD(G)$ and $DD(G')$ are multisets that contain the DAGs extracted from G and G' , respectively, and k_{DAG} is a kernel between DAGs

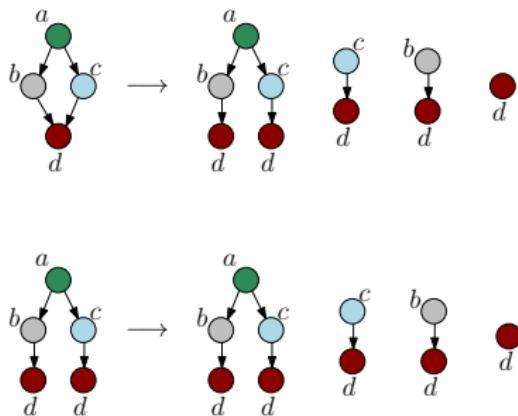
[Da San Martino et al., SDM'12]

Ordered Decomposition DAGs Kernel

DAGs are unordered (i.e. the set of neighbours of each node is unordered)

There is a vast literature on kernels for ordered trees. Hence, the kernel:

- transforms the unordered DAGs to ordered DAGs (based on node labels, outdegrees of nodes, etc.)
- projects subdags to a tree space (see Figure below)
- applies a kernel for ordered trees



The kernel between two DAGs is computed as follows:

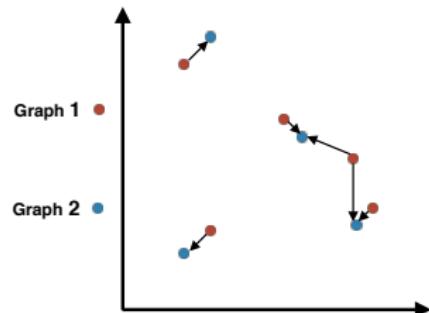
$$k_{DAG}(D, D') = \sum_{v \in V_D} \sum_{v' \in V_{D'}} k_{tree}(\text{root}(v), \text{root}(v'))$$

$V_D, V_{D'}$: sets of vertices of D and D'
 k_{tree} : kernel between ordered trees

Assignment Kernels

Assignment Kernels

- Another design paradigm for developing kernels
- Only a few instances in the literature
- They compute a matching between substructures of one object and substructures of a second object such that the overall similarity of the two objects is maximized
- Such a matching can reveal structural correspondences between the two objects



Pyramid Match Kernel

Embed all vertices in the d -dimensional vector space \mathbb{R}^d as follows

- compute the eigendecomposition of the adjacency matrix
- use the eigenvectors of the d largest in magnitude eigenvalues

Such embeddings capture **global** properties of graphs

Example: eigenvector corresponding to greatest eigenvalue contains eigenvector centrality scores of vertices → global property

After embedding: each vertex is a point in the d -dimensional unit hypercube

Then, use pyramid match kernel, a kernel function over unordered feature sets:

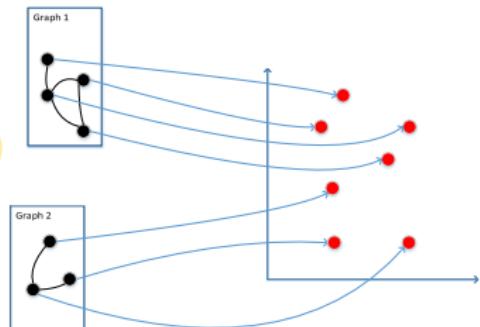
- Each feature set is mapped to a multiresolution histogram
- The histogram pyramids are then compared using a weighted histogram intersection computation

[Nikolentzos et al., AAAI'17]

Node Embeddings

Node embeddings: represent nodes as points in a vector space

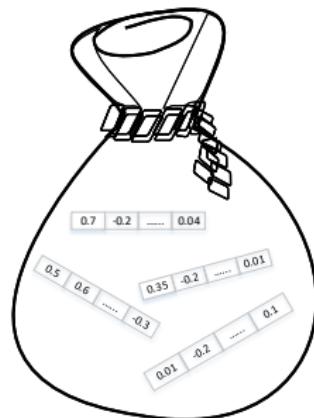
- Generate embeddings using eigenvectors of adjacency matrix $A = U \Lambda U^\top$
 - i^{th} row u_i of U corresponds to embedding of vertex v_i
- Such embeddings capture global properties of graphs



Bag-of-vectors Representation

Graphs represented as **bags-of-vectors**:

- A graph is represented as a set of vectors:
 $\{u_1, \dots, u_n\}$
- Each vector $u_i \in \mathbb{R}^d$ represents the embedding of the i^{th} node in the d -dimensional space
- This is a natural representation
 \hookrightarrow There is no canonical ordering for the nodes of a graph



Pyramid Match Graph Kernel

The Pyramid Match Graph Kernel

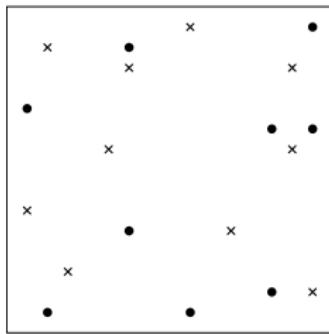
- partitions feature space into cells
- at level $l \rightarrow 2^l$ cells along each dimension

Number of nodes (i.e. embeddings) that match at l :

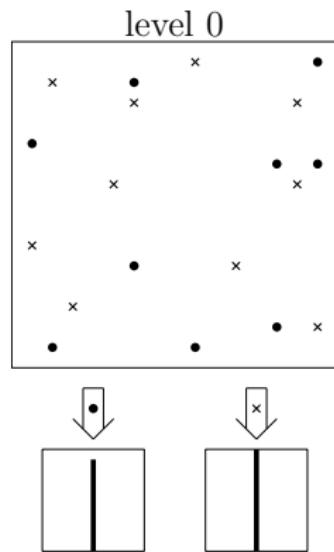
$$I(H_{G_1}^l, H_{G_2}^l) = \sum_{i=1}^{2^l d} \min(H_{G_1}^l(i), H_{G_2}^l(i))$$

where $H_G^l(i)$ is the number of nodes of G that lie in the i^{th} cell

Example

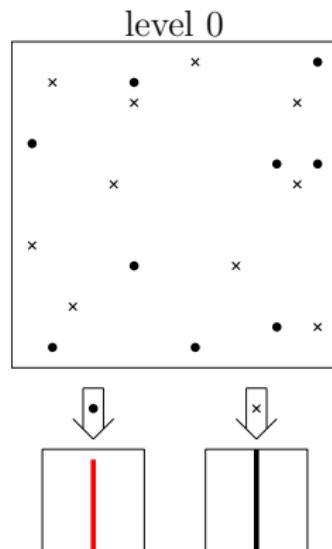


Example



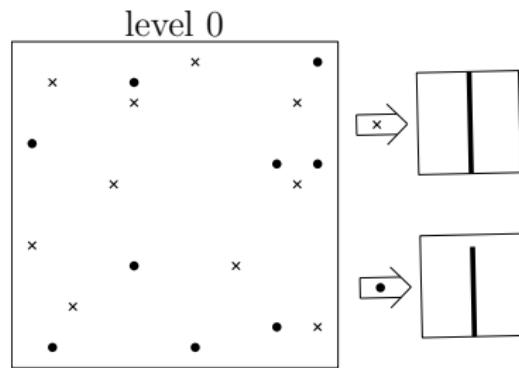
Example

$$I(H_{G_1}^0, H_{G_2}^0) = 9 + \dots$$



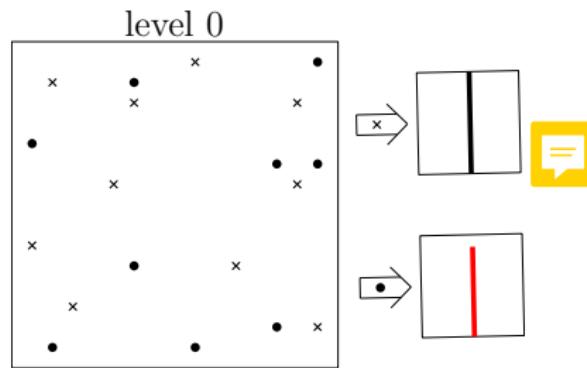
Example

$$I(H_{G_1}^0, H_{G_2}^0) = 9 + \dots$$



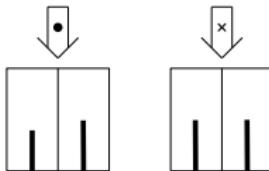
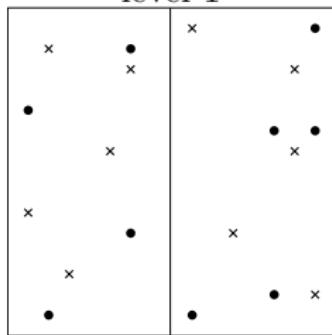
Example

$$I(H_{G_1}^0, H_{G_2}^0) = 9 + 9 = 18$$



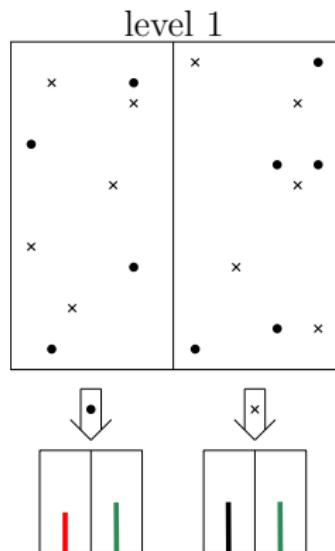
Example

level 1



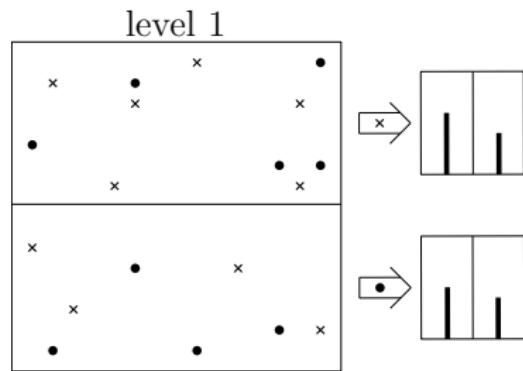
Example

$$I(H_{G_1}^1, H_{G_2}^1) = (5 + 4) + \dots$$



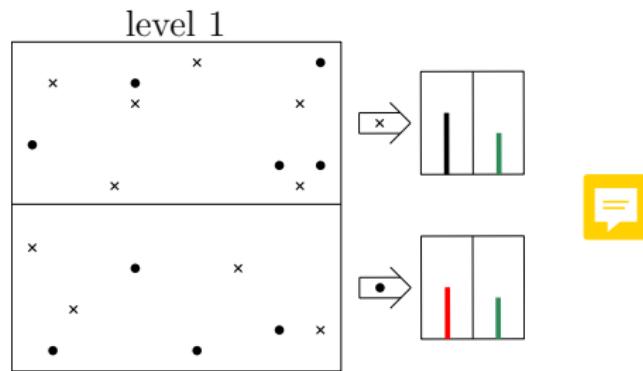
Example

$$I(H_{G_1}^1, H_{G_2}^1) = (5 + 4) + \dots$$

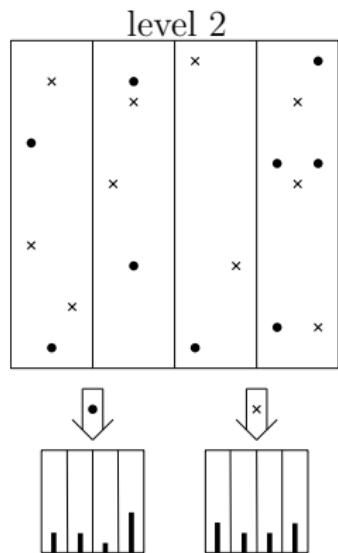


Example

$$I(H_{G_1}^1, H_{G_2}^1) = (5 + 4) + (5 + 4) = 18$$

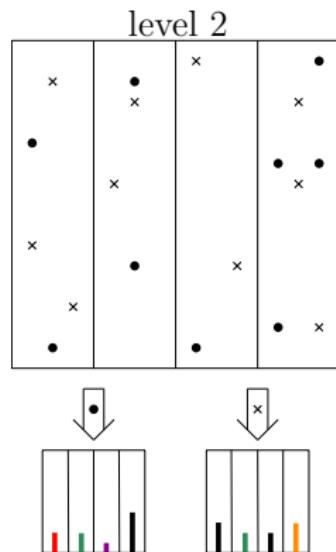


Example



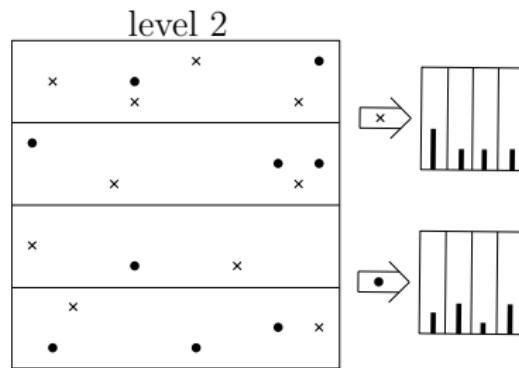
Example

$$I(H_{G_1}^2, H_{G_2}^2) = (2 + 2 + 1 + 3) + \dots$$



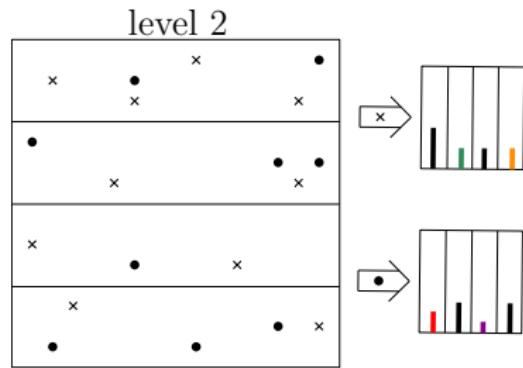
Example

$$I(H_{G_1}^2, H_{G_2}^2) = (2 + 2 + 1 + 3) + \dots$$



Example

$$I(H_{G_1}^2, H_{G_2}^2) = (2 + 2 + 1 + 3) + (2 + 2 + 1 + 2) = 15$$



Pyramid Match Graph Kernel

PM takes a weighted sum of the matches that occur at each level (levels 0 to L):

$$\begin{aligned} k_{\Delta}(G_1, G_2) &= I(H_{G_1}^L, H_{G_2}^L) + \sum_{l=0}^{L-1} \frac{1}{2^{L-l}} (I(H_{G_1}^l, H_{G_2}^l) - I(H_{G_1}^{l+1}, H_{G_2}^{l+1})) \\ &= 15 + \frac{1}{2}(18 - 15) + \frac{1}{4}(18 - 18) = 16.5 \end{aligned}$$

- Matches within lower levels weighted less
- Only new matches are taken into account

Complexity: $\mathcal{O}(dnL)$

Optimal Assignment Kernel

- $\{x_1, \dots, x_n\}$ are substructures of G , e.g., nodes
- $\{x'_1, \dots, x'_{n'}\}$ are substructures of G' , e.g., nodes
- κ is a non-negative kernel comparing substructures
- π is a permutation of the integers $\{1, \dots, \min(n, n')\}$
- Then, the optimal assignment kernel is defined as follows:

$$k(G, G') = \begin{cases} \max_{\pi} \sum_{i=1}^n \kappa(x_i, x'_{\pi(i)}), & \text{if } n' > n \\ \max_{\pi} \sum_{j=1}^{n'} \kappa(x_{\pi(j)}, x'_j), & \text{otherwise} \end{cases}$$

[Frohlich et al., ICML'05]

Optimal Assignment Kernel

- $\{x_1, \dots, x_n\}$ are substructures of G , e.g., nodes
- $\{x'_1, \dots, x'_{n'}\}$ are substructures of G' , e.g., nodes
- κ is a non-negative kernel comparing substructures
- π is a permutation of the integers $\{1, \dots, \min(n, n')\}$
- Then, the optimal assignment kernel is defined as follows:

$$k(G, G') = \begin{cases} \max_{\pi} \sum_{i=1}^n \kappa(x_i, x'_{\pi(i)}), & \text{if } n' > n \\ \max_{\pi} \sum_{j=1}^{n'} \kappa(x_{\pi(j)}, x'_j), & \text{otherwise} \end{cases}$$

- However, **not** positive semidefinite in general

[Vert. arXiv:0801.4061]

Valid Optimal Assignment Kernels

- Let \mathcal{X} be a set, and $[\mathcal{X}]^n$ denote the set of all n -element subsets of \mathcal{X}
- Let also $X, X' \in [\mathcal{X}]^n$ for $n \in \mathbb{N}$, and $\mathfrak{B}(X, X')$ denote the set of all bijections between X and X'
- The optimal assignment kernel on $[\mathcal{X}]^n$ is defined as

$$K_{\mathfrak{B}}^k(X, X') = \max_{B \in \mathfrak{B}(X, X')} \sum_{(x, x') \in B} k(x, x')$$

where k is a kernel between the elements of X and X'

- The above function $K_{\mathfrak{B}}(\mathcal{X}, \mathcal{X}')$ is a valid kernel only if the base kernel k is strong

Definition (Strong Kernel)

A function $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}_{\geq 0}$ is called strong kernel if $k(x, y) \geq \min\{k(x, z), k(z, y)\}$ for all $x, y, z \in \mathcal{X}$.

Strong kernels are equivalent to kernels obtained from a *hierarchy* defined on set

\mathcal{X}

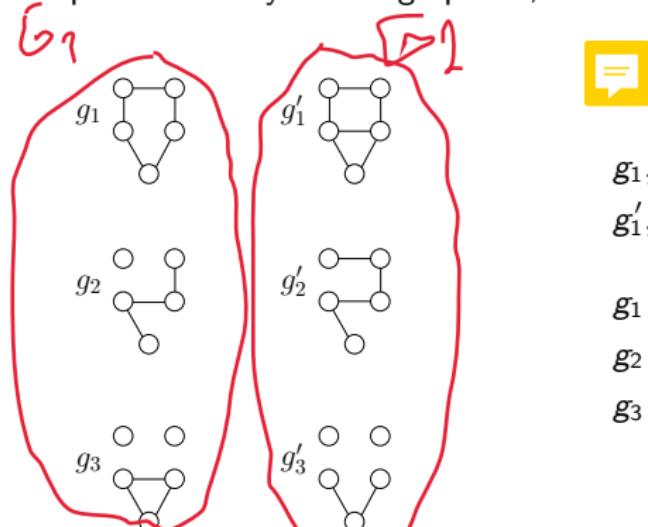
Frameworks

Diagonal Dominance Problem

Diagonal dominance problem of kernels that compare specific substructures of graphs:

- Very large feature space, hence, unlikely that two graphs will contain similar substructures
- However, substructures (i. e. features) often **related** to each other
- Kernel value between pairs of graphs \ll kernel value between a graph and itself

For example, when the features correspond to large graphlets (e.g., $k \geq 5$), two graphs may be composed of many similar graphlets, but not any identical



g_1, g_2, g_3 extracted from G

g'_1, g'_2, g'_3 extracted from G'

g_1 nearly isomorphic to g'_1

g_2 nearly isomorphic to g'_2

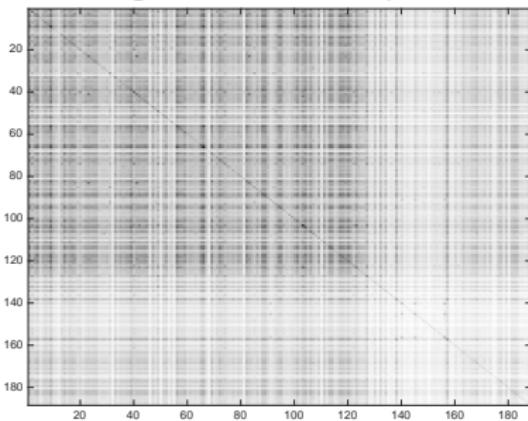
g_3 nearly isomorphic to g'_3

Diagonal Dominance Problem

Diagonal dominance problem of kernels that compare specific substructures of graphs:

- Very large feature space, hence, unlikely that two graphs will contain similar substructures
- However, substructures (i. e. features) often **related** to each other
- Kernel value between pairs of graphs \ll kernel value between a graph and itself

This leads to the diagonal dominance problem



The resulting kernel matrix is close to the identity matrix

A Structural Smoothing Framework

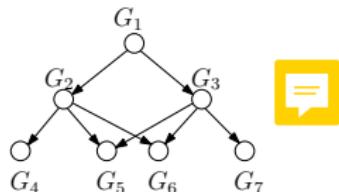
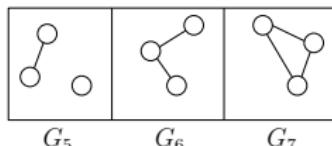
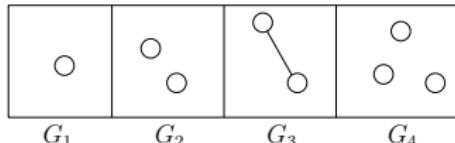
To deal with diagonal dominance, it applies *smoothing*

First construct a Directed Acyclic Graph (DAG):

- each vertex corresponds to a substructure
- for each substructure s of size k determine all possible substructures of size $k - 1$ that s can be reduced into
- these correspond to the parents of s
- draw a weighted directed edge from each parent to its children vertices

DAG provides a topological ordering of the vertices

- all descendants of a given substructure at depth $k - 1$ are at depth k



DAG for graphlets of size $k \leq 3$

[Yanardag and Vishwanathan, NIPS'15]

A Structural Smoothing Framework

The structural smoothing for a substructure s at level k is defined as:

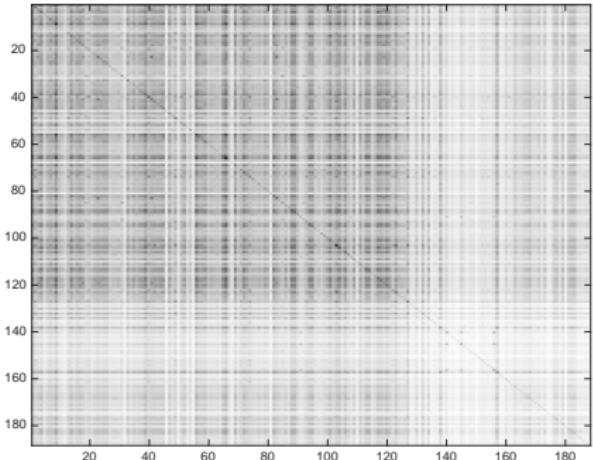
$$P_{SS}^k(s) = \frac{\max(c_s - d, 0)}{m} + \frac{dm_d}{m} \sum_{p \in \mathcal{P}_s} P_{SS}^{k-1}(p) \frac{w_{ps}}{\sum_{c \in \mathcal{C}_p} w_{pc}}$$

where

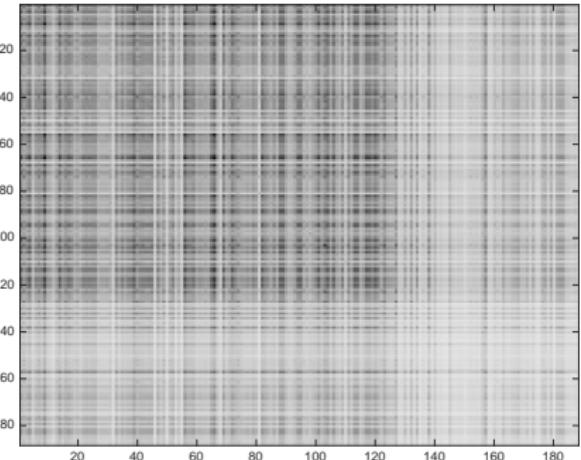
- c_s denotes the number of times substructure s appears in the graph
- $m = \sum_i c_i$ denotes the total number of substructures present in the graph
- $d > 0$ is a discount factor
- $m_d := |\{i : c_i > d\}|$ is the number of substructures whose counts are larger than d
- w_{ij} denotes the weight of the edge connecting vertex i to vertex j
- \mathcal{P}_s denotes the parents of vertex s
- \mathcal{C}_p the children of vertex p

Even if the graph does not contain a substructure s ($c_s = 0$), its value in the feature vector can be **greater** than 0 ($P_{SS}(s) > 0$)

Example



Kernel matrix before
smoothing



Kernel matrix after
smoothing

Deep Graph Kernels

To deal with diagonal dominance, the deep graph kernels framework computes the kernel as follows:

$$k(G, G') = \phi(G)^\top M \phi(G')$$

M : a positive semidefinite matrix that encodes the relationships between substructures
Each component of $\phi(G), \phi(G')$ corresponds to a substructure (e.g., the complete graphlet of size 5)

Matrix M is learned using techniques inspired from the field of natural language processing:

- An embedding for each substructure is generated using the CBOW or Skip-gram model
- Then M corresponds to the inner products of these embeddings

However, unlike words in documents, substructures of graphs do not have a *linear co-occurrence relationship*

Such co-occurrence relationships are manually defined for 3 kernels:

- (1) the Weisfeiler-Lehman subtree kernel
- (2) the graphlet kernel
- (3) the shortest path kernel

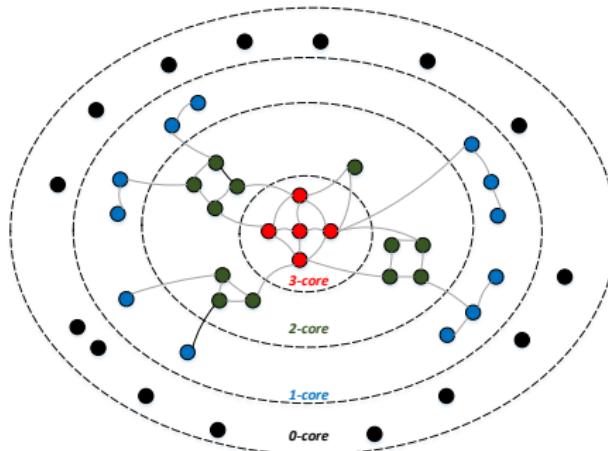
[Yanardag and Vishwanathan, KDD'15]

A Degeneracy Framework for Graph Comparison

Definition (k -core)

The k -core of a graph is defined as a maximal subgraph in which every vertex is connected to at least k other vertices within that subgraph

A k -core decomposition of a graph consists of finding the set of all k -cores



The set of all k -cores forms a nested sequence of subgraphs

The degeneracy $\delta^*(G)$ is defined as the maximum k for which graph G contains a non-empty k -core subgraph

[Nikolentzos et al., IJCAI'18]

Degeneracy Framework for Graph Comparison

Idea: use the nested sequence of subgraphs generated by k -core decomposition to capture structure at multiple different scales

The core variant of the base kernel k is defined as:

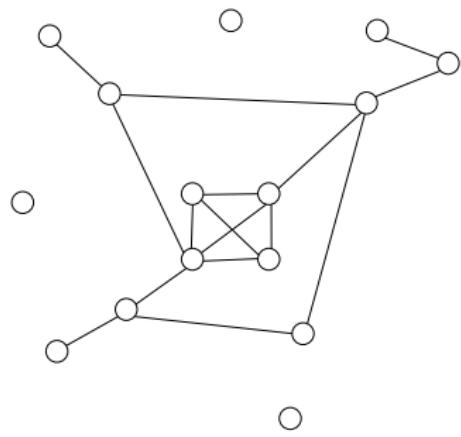
$$k_c(G, G') = k(C_0, C'_0) + k(C_1, C'_1) + \dots + k(C_{\delta_{min}^*}, C'_{\delta_{min}^*})$$

where δ_{min}^* is the minimum of the degeneracies of the two graphs, and $C_0, C_1, \dots, C_{\delta_{min}^*}$ and $C'_0, C'_1, \dots, C'_{\delta_{min}^*}$ are the 0-core, 1-core, ..., δ_{min}^* -core subgraphs of G and G' , respectively

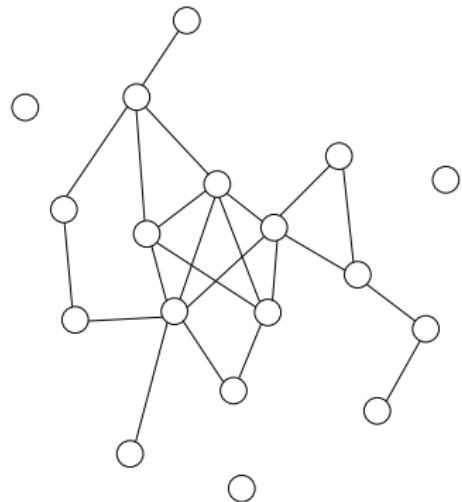
The degeneracy framework can:

- increase the expressive power of existing algorithms
- be applied to any algorithm that compares graphs

Example

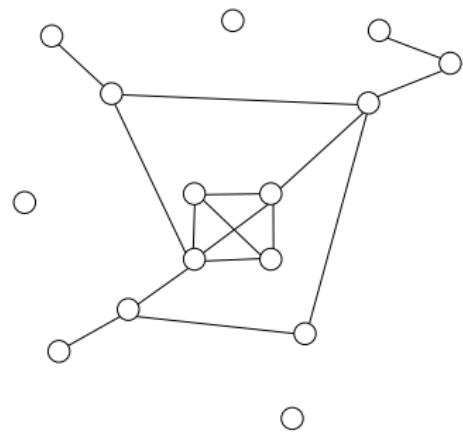


G

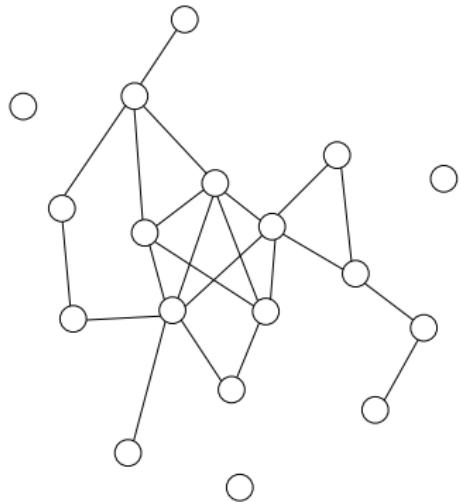


G'

Example



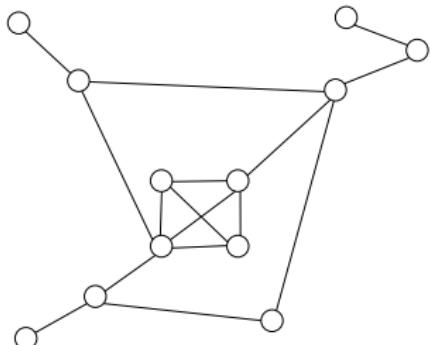
C_0



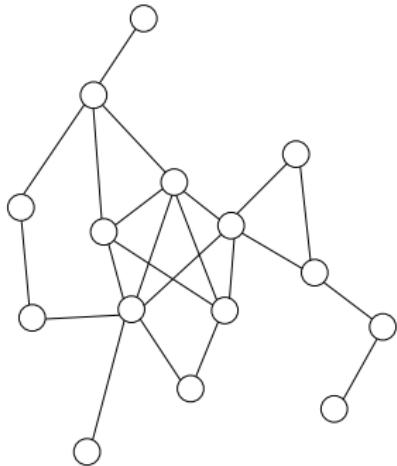
C'_0

$$k_c(G, G') = k(C_0, C'_0)$$

Example



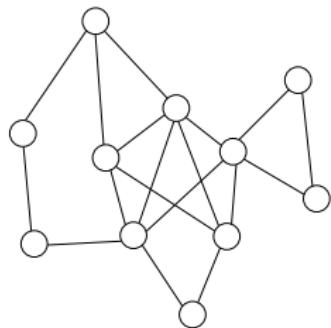
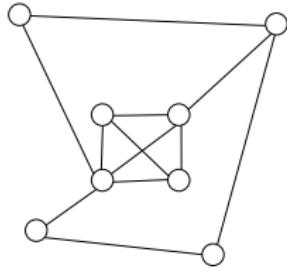
C_1



C'_1

$$k_c(G, G') = k(C_0, C'_0) + k(C_1, C'_1)$$

Example



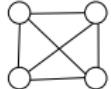
C_2



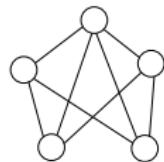
C'_2

$$k_c(G, G') = k(C_0, C'_0) + k(C_1, C'_1) + k(C_2, C'_2)$$

Example



C_3



C'_3

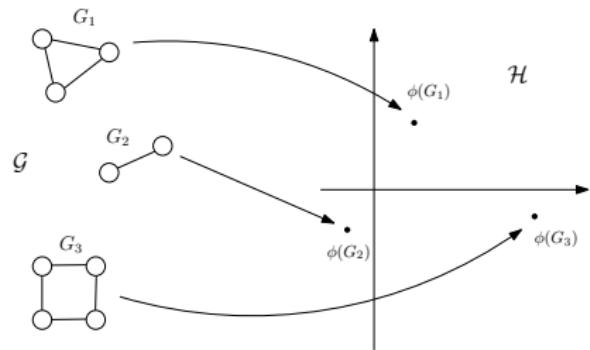


$$k_c(G, G') = k(C_0, C'_0) + k(C_1, C'_1) + k(C_2, C'_2) + k(C_3, C'_3)$$

Successive Embeddings

Graph kernels compute implicitly the inner product between the representations of input graphs in \mathcal{H}

- Equivalent to computing the linear kernel on feature space \mathcal{H}
- Linear kernel limits expressiveness of derived representations



Idea: Obtain complex kernels by stacking simpler kernels on top of one another

[Nikolentzos et al., CIKM'18]

Successive Embeddings

Embedding 1: Embed graphs in a Hilbert space \mathcal{H}_1 using a graph kernel k

Embedding 2: Embed emerging representations $\phi(G), \phi(G')$ into another Hilbert space \mathcal{H}_2 using kernels for vector data:

① *Polynomial kernel:* $k_P(\phi(G), \phi(G')) = (\langle \phi(G), \phi(G') \rangle)^d, \quad d \in \mathbb{N}$

② *Gaussian kernel:* $k_G(\phi(G), \phi(G')) = \exp\left(-\frac{\|\phi(G) - \phi(G')\|^2}{2\sigma^2}\right), \quad \sigma > 0$

Problem: Usually $\phi(G)$ and $\phi(G')$ not computed explicitly. How to apply

Embedding 2?

↪ Use an implicit computation scheme

The two kernels for vector data can be computed as:

① *Polynomial kernel:*

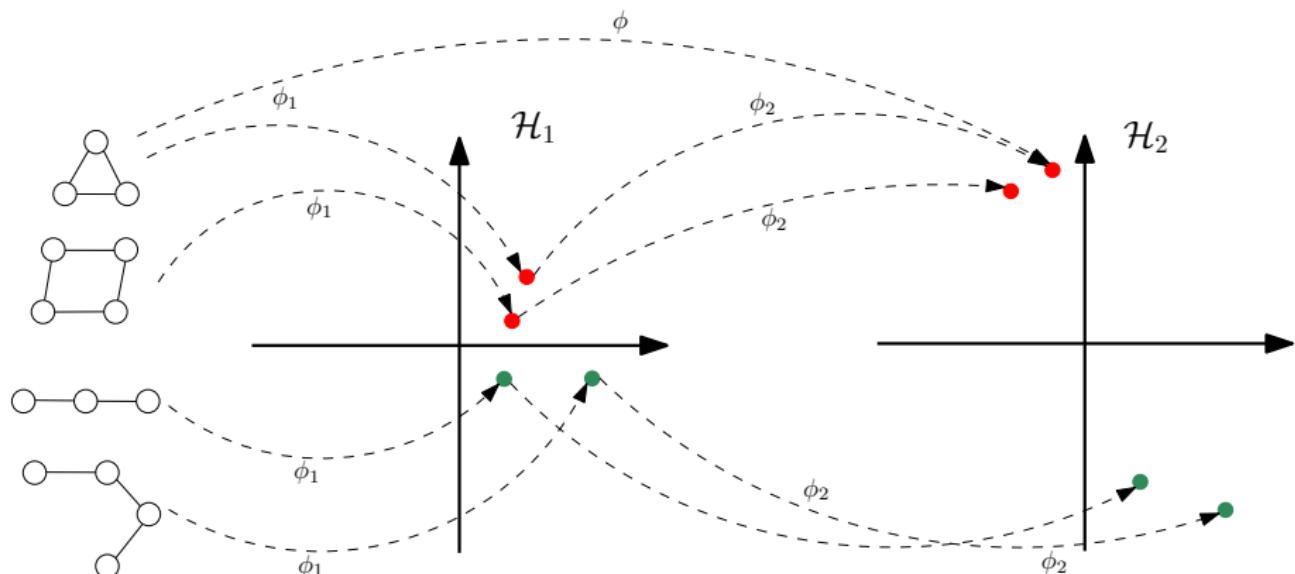
$$k_P(\phi(G), \phi(G')) = (\langle \phi(G), \phi(G') \rangle)^d = (k(G, G'))^d, \quad d \in \mathbb{N}$$

② *Gaussian kernel:*

$$k_G(\phi(G), \phi(G')) = \exp\left(-\frac{k(G, G) - 2k(G, G') + k(G', G')}{2\sigma^2}\right), \quad \sigma > 0$$

Successive Embeddings Example

- Figure below illustrates a sequence of two embeddings
- Separation of the data points associated with the two classes progressively increased



Applications of Graph Kernels

Applications

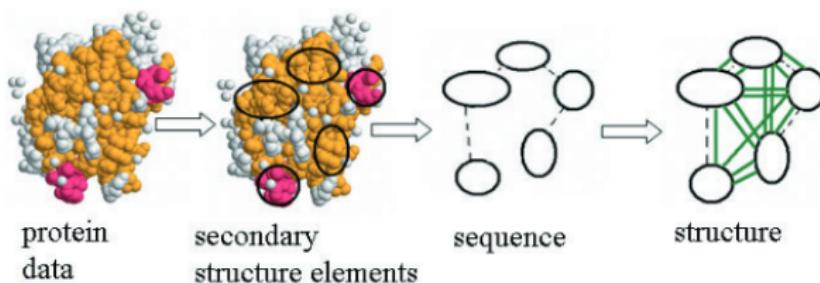
- Bioinformatics [Borgwardt et al., Bioinformatics 21(suppl_1); Borgwardt et al., PSB'07; Sato et al., BMC bioinformatics 9(1)]
- Chemoinformatics [Swamidass et al., Bioinformatics 21(suppl_1); Ralaivola et al., Neural Networks 18(8); Mahé et al., JCIM 45(4); Ceroni et al., Bioinformatics 23(16); Mahé and Vert, Machine Learning 75(1)]
- Computer Vision [Harchaoui and Bach, CVPR'07; Bach, ICML'08; Wang and Sahbi. CVPR'13; Stumm et al., CVPR'16]
- Cybersecurity [Anderson et al., JCV 7(4); Gascon et al., AISeC'13; Narayanan et al., IJCNN'16]
- Natural Language Processing [Glavas and Snajder, ACL'13; Bleik et al., TCBB 10(5); Nikolentzos et al., EMNLP'17]
- Social Networks [Yanardag and Vishwanathan, KDD'15]

:

Protein Function Prediction

For each protein, create a graph that contains information about its

- structure
- sequence
- chemical properties

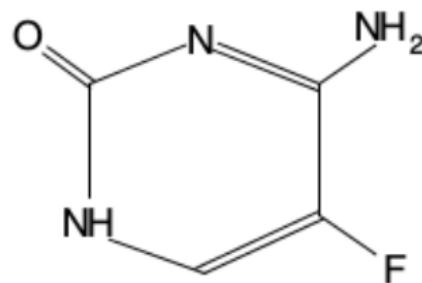


Perform **graph classification** to predict the function of proteins

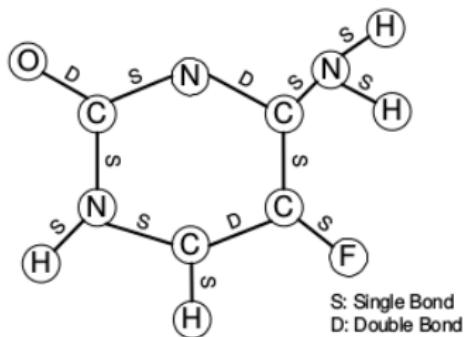
| Kernel type | Accuracy |
|--------------------------------|----------|
| Vector kernel | 76.86 |
| Optimized vector kernel | 80.17 |
| Graph kernel | 77.30 |
| Graph kernel without structure | 72.33 |
| Graph kernel with global info | 84.04 |
| DALI classifier | 75.07 |

Chemical Compound Classification

Represent each chemical compound as a graph



⇒



S: Single Bond
D: Double Bond

Perform **graph classification** to predict if a chemical compound displays the desired behavior against the specific biomolecular target or not

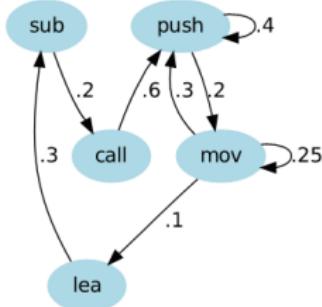
| | graph kernels | | | | | | | |
|---------|---------------|-------|---------|---------|-------|--------|---------------|--|
| Lin.Reg | DT | NN | Progoll | Progol2 | Sebag | Kramer | graph kernels | |
| 89.3% | 88.3% | 89.4% | 81.4% | 87.8% | 93.3% | 95.7% | 91.2% | |

[Mahé et al., JCIM 45(4)]

Malware Detection

Given a computer program, create its control flow graph

| | |
|------|----------------------|
| call | [ebp+0x8] |
| push | 0x70 |
| push | 0x010012F8 |
| call | 0x01006170 |
| push | 0x010061C0 |
| mov | eax, fs:[0x00000000] |
| push | eax |
| mov | fs:[], esp |
| mov | eax, [esp+0x10] |
| mov | [esp+0x10], ebp |
| lea | ebp, [esp+0x10] |
| sub | esp, eax |
| ... | ... |



Perform **graph classification** to predict if there is malicious code inside the program or not

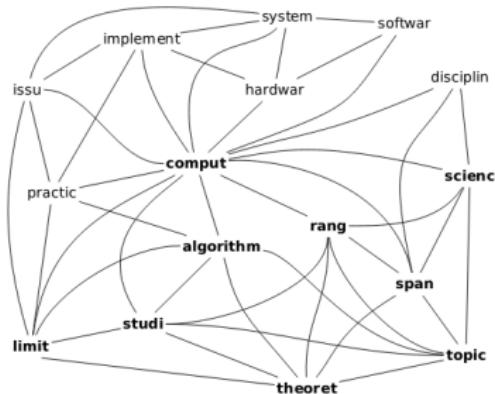
| Method | Accuracy (%) |
|--|---------------|
| Gaussian kernel | 99.09 |
| Spectral kernel | 96.36 |
| Combined kernel | 100.00 |
| n-gram ($n = 4$, $L = 1,000$, SVM = 2-poly) | 94.55 |
| n-gram ($n = 4$, $L = 2,500$, SVM = Gauss) | 93.64 |
| n-gram ($n = 6$, $L = 2,500$, SVM = 2-poly) | 92.73 |
| n-gram ($n = 3$, $L = 1,000$, SVM = 2-poly) | 89.09 |
| n-gram ($n = 2$, $L = 500$, 3-NN) | 88.18 |

Graph-Of-Words

Each document is represented as a graph
 $G = (V, E)$ consisting of a set V of vertices and a set E of edges between them

- vertices → unique terms
- edges → co-occurrences within a fixed-size sliding window
- no edge weight
- no edge direction

As a discipline, computer science spans a range of topics from theoretical studies of algorithms and the limits of computation to the practical issues of implementing computing systems in hardware and software.



Graph representation more flexible than n -grams. Takes into account

- word inversion
- subset matching
- e.g., “article about news” vs. “news article”

Custom Shortest Path Kernel

Transforms the original graphs into shortest-paths graphs

↪ Edges correspond to shortest paths of length at most d in original graph

Given the SP-transformed graphs $C_1 = (V_1, E_1)$ and $C_2 = (V_2, E_2)$ of G_1 and G_2 , the shortest path kernel is defined as:

$$k(G_1, G_2) = \frac{\sum_{v_1 \in V_1, v_2 \in V_2} k_{node}(v_1, v_2) + \sum_{e_1 \in E_1, e_2 \in E_2} k_{walk}^{(1)}(e_1, e_2)}{norm}$$

where k_{node} is a kernel for comparing two vertices, $k_{walk}^{(1)}$ a kernel on edge walks of length 1 and $norm$ a normalization factor. Specifically:

$$k_{node}(v_1, v_2) = \begin{cases} 1 & \text{if } \ell(v_1) = \ell(v_2), \\ 0 & \text{otherwise} \end{cases}$$

$$k_{walk}^{(1)}(e_1, e_2) = k_{node}(u_1, u_2) k_{edge}(e_1, e_2) k_{node}(v_1, v_2)$$

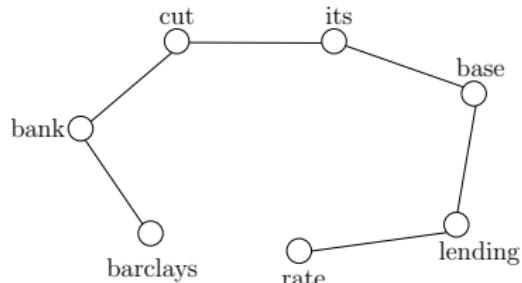
$$k_{edge}(e_1, e_2) = \begin{cases} \ell(e_1) \ell(e_2) & \text{if } e_1 \in E_1 \wedge e_2 \in E_2, \\ 0 & \text{otherwise} \end{cases}$$

[Nikolentzos et al., EMNLP'17]

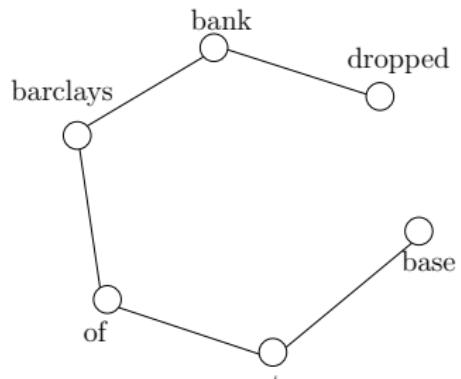
Example

d_1 : “barclays bank cut its base lending rate”

d_2 : “base rate of barclays bank dropped”



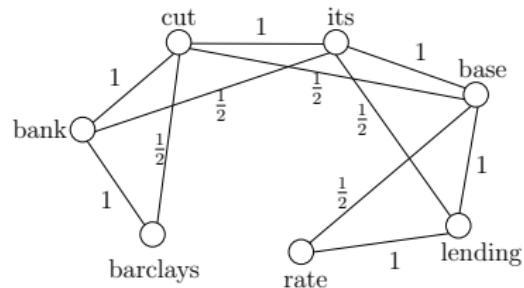
G_1



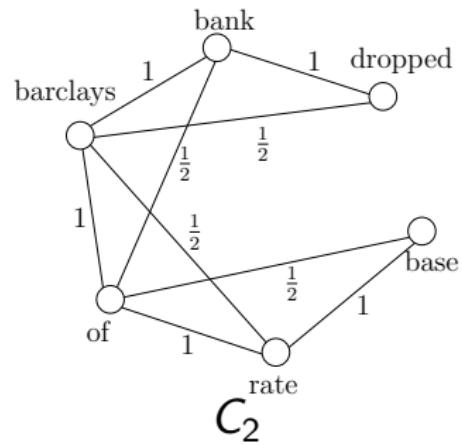
G_2

Example

SP-transformation ($d = 2$)

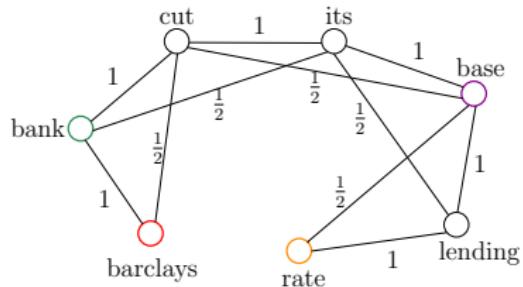


C_1

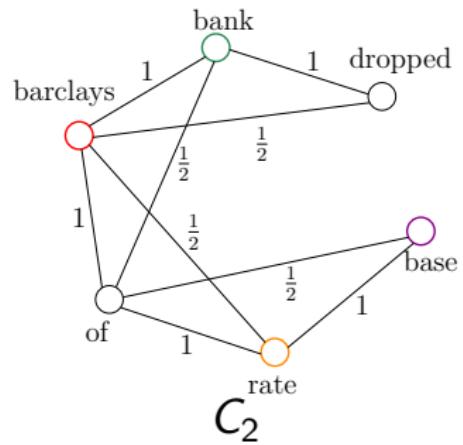


Example

$$\sum_{v_1 \in V_1, v_2 \in V_2} k_{node}(v_1, v_2) = 4$$

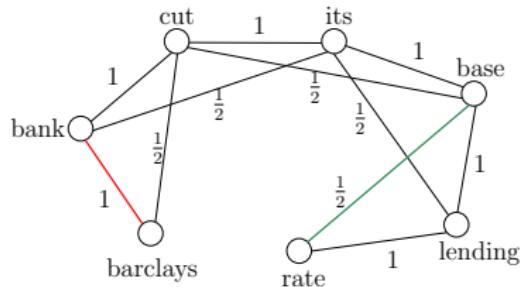


C_1

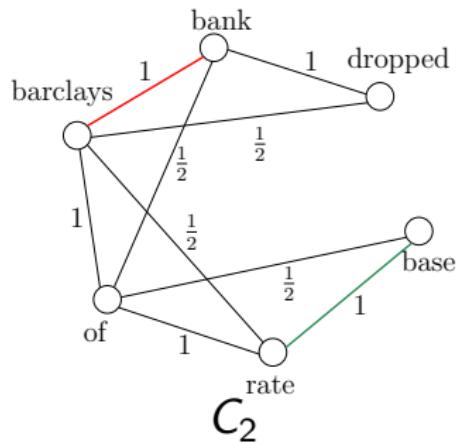


Example

$$\sum_{e_1 \in E_1, e_2 \in E_2} k_{\text{walk}}^{(1)}(e_1, e_2) = 1 + \frac{1}{2} = \frac{3}{2}$$



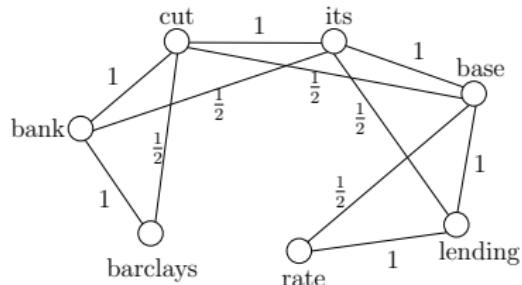
C_1



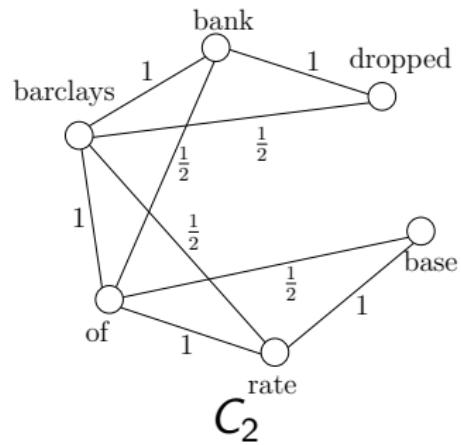
Example

$norm = 13.07$

$$k(G_1, G_2) = \frac{4 + \frac{3}{2}}{13.07} = 0.42$$



C_1

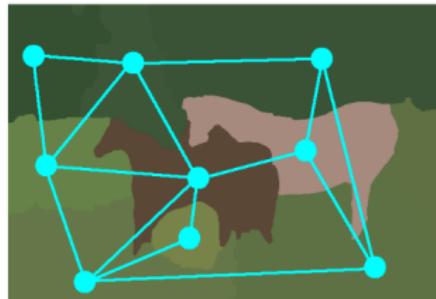


Text Categorization

| Method | Dataset | WebKB | | News | | Subjectivity | | Amazon | | Polarity | |
|-------------|-----------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|
| | | Acc | F1 |
| Dot product | $n = 1$ | 90.26 | 89.23 | 81.10 | 77.64 | 89.92 | 89.92 | 91.88 | 91.88 | 76.27 | 76.26 |
| | $n = 2$ | 90.47 | 89.50 | 80.91 | 77.32 | 91.01 | 91.01 | 92.00 | 92.02 | 77.46 | 77.45 |
| | $n = 3$ | 90.26 | 89.17 | 80.72 | 77.10 | 90.90 | 90.90 | 91.81 | 91.85 | 77.41 | 77.40 |
| | $n = 4$ | 89.40 | 88.13 | 80.31 | 76.51 | 90.39 | 90.39 | 91.31 | 91.33 | 77.19 | 77.18 |
| Cosine | $n = 1$ | 92.48 | 91.88 | 81.17 | 77.66 | 90.03 | 90.02 | 94.00 | 94.00 | 76.70 | 76.69 |
| | $n = 2$ | 93.05 | 92.75 | 81.49 | 77.97 | 90.94 | 90.94 | 94.13 | 94.13 | 77.56 | 77.56 |
| | $n = 3$ | 92.98 | 92.59 | 80.97 | 77.38 | 90.99 | 90.99 | 94.19 | 94.18 | 77.65 | 77.65 |
| | $n = 4$ | 92.48 | 92.08 | 80.76 | 77.09 | 90.76 | 90.75 | 94.13 | 94.13 | 77.53 | 77.53 |
| Tanimoto | $n = 1$ | 90.62 | 89.83 | 81.55 | 78.15 | 90.94 | 90.93 | 92.25 | 92.26 | 77.49 | 77.48 |
| | $n = 2$ | 90.40 | 89.45 | 80.75 | 77.00 | 90.61 | 90.60 | 91.81 | 91.85 | 77.35 | 77.35 |
| | $n = 3$ | 92.41 | 91.80 | 79.80 | 75.75 | 90.21 | 90.20 | 93.44 | 93.47 | 76.48 | 76.48 |
| | $n = 4$ | 91.76 | 90.84 | 78.99 | 74.83 | 89.53 | 89.52 | 93.00 | 93.00 | 75.86 | 75.86 |
| DCNN | | 89.18 | 87.99 | 79.91 | 76.15 | 90.26 | 90.26 | 91.81 | 91.81 | 73.26 | 73.26 |
| CNN | static,rand | > 1 day | | 77.57 | 73.37 | 87.16 | 87.15 | 88.81 | 88.82 | 71.50 | 71.50 |
| | non-static,rand | > 1 day | | 81.13 | 77.49 | 89.61 | 89.60 | 93.56 | 93.56 | 76.54 | 76.53 |
| SPGK | $d = 1$ | 93.27 | 92.78 | 81.04 | 77.49 | 91.48 | 91.48 | 94.00 | 94.01 | 77.76 | 77.75 |
| | $d = 2$ | 93.70 | 93.36 | 80.89 | 77.29 | 91.46 | 91.46 | 94.13 | 94.13 | 77.89 | 77.88 |
| | $d = 3$ | 92.91 | 92.33 | 80.78 | 77.03 | 91.37 | 91.37 | 94.44 | 94.44 | 77.61 | 77.60 |
| | $d = 4$ | 92.91 | 92.23 | 80.97 | 77.30 | 91.18 | 91.18 | 94.63 | 94.63 | 77.80 | 77.80 |

Image Classification

Represent each image as a graph based on its segmentation mosaic



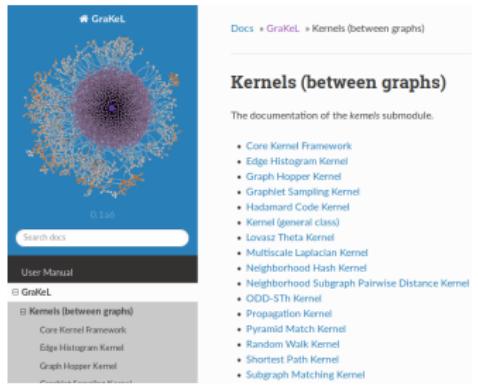
Perform **graph classification** to categorize images

| | H | W | TW | wTW | M |
|---------|--------|-------|-------|-------|-------|
| Coil100 | 1.2% | 0.8% | 0.0% | 0.0% | 0.0% |
| Corel14 | 10.36% | 8.52% | 7.24% | 6.12% | 5.38% |

[Harchaoui and Bach, CVPR'07]

Experimental Evaluation

- Python library for graph kernels
- Contains implementations of a large number of graph kernels
- Compatible with scikit-learn
- Project repository:
<https://github.com/ysig/GraKeL>



[Siglidis et al., arXiv:1806.02193]

Evaluation

Standard datasets from graph classification containing:

- unlabeled graphs
- node-labeled graphs
- node-attributed graphs

Classification using:

- SVM → precompute kernel matrix
- Hyperparameters of both SVM (i.e. C) and graph kernels optimized on training set using cross-validation

Perform 10 **times** 10-fold cross validation and report:

- *Average accuracy* over the 10 repetitions
- Standard deviation over the 10 repetitions

Graph Classification (Node-Labeled Graphs)

| KERNELS | DATASETS | | | |
|---|---------------------|---------------------|---------------------|---------------------|
| | MUTAG | ENZYMES | NCI1 | PTC-MR |
| VERTEX HISTOGRAM | 71.87 (\pm 1.83) | 16.87 (\pm 1.56) | 56.09 (\pm 0.35) | 58.09 (\pm 0.62) |
| RANDOM WALK | 82.24 (\pm 2.87) | 12.90 (\pm 1.42) | TIMEOUT | 51.26 (\pm 2.30) |
| SHORTEST PATH | 82.54 (\pm 1.00) | 40.13 (\pm 1.34) | 72.25 (\pm 0.28) | 59.26 (\pm 2.34) |
| WL SUBTREE | 84.00 (\pm 1.25) | 53.15 (\pm 1.22) | 85.03 (\pm 0.20) | 63.28 (\pm 1.34) |
| WL SHORTEST PATH | 82.29 (\pm 1.93) | 28.23 (\pm 1.00) | 61.43 (\pm 0.32) | 55.51 (\pm 1.68) |
| WL PYRAMID MATCH | 88.60 (\pm 0.95) | 57.72 (\pm 0.84) | 85.31 (\pm 0.42) | 64.52 (\pm 1.36) |
| NEIGHBORHOOD HASH | 87.74 (\pm 1.17) | 43.43 (\pm 1.45) | 74.81 (\pm 0.37) | 60.50 (\pm 2.10) |
| NEIGHBORHOOD SUBGRAPH PAIRWISE DISTANCE | 82.46 (\pm 1.55) | 41.97 (\pm 1.66) | 74.36 (\pm 0.31) | 60.04 (\pm 1.15) |
| ORDERED DAGs DECOMPOSITION | 79.01 (\pm 2.04) | 31.87 (\pm 1.35) | 75.03 (\pm 0.45) | 59.08 (\pm 1.85) |
| PYRAMID MATCH | 84.72 (\pm 1.67) | 42.67 (\pm 1.78) | 73.11 (\pm 0.49) | 57.99 (\pm 2.45) |
| GRAPHHOPPER | 82.11 (\pm 2.13) | 36.47 (\pm 2.13) | 71.36 (\pm 0.13) | 55.64 (\pm 2.03) |
| SUBGRAPH MATCHING | 84.04 (\pm 1.55) | 35.68 (\pm 0.80) | TIMEOUT | 57.91 (\pm 1.73) |
| PROPAGATION | 77.23 (\pm 1.22) | 44.48 (\pm 1.63) | 82.12 (\pm 0.22) | 59.30 (\pm 1.24) |
| MULTISCALE LAPLACIAN | 86.11 (\pm 1.60) | 53.08 (\pm 1.53) | 79.40 (\pm 0.47) | 59.95 (\pm 1.71) |
| CORE WL | 85.90 (\pm 1.44) | 52.37 (\pm 1.29) | 85.12 (\pm 0.21) | 63.03 (\pm 1.67) |
| CORE SHORTEST PATH | 85.13 (\pm 2.46) | 41.55 (\pm 1.66) | 73.87 (\pm 0.19) | 58.21 (\pm 1.87) |

| KERNELS | DATASETS | | | AVG. RANK |
|---|---------------------|---------------------|---------------------|--------------|
| | D&D | PROTEINS | AIDS | |
| VERTEX HISTOGRAM | 74.83 (\pm 0.40) | 70.93 (\pm 0.28) | 79.78 (\pm 0.13) | 13.7 |
| RANDOM WALK | OUT-OF-MEM | 69.31 (\pm 0.29) | 79.52 (\pm 0.58) | 15.0 |
| SHORTEST PATH | 78.93 (\pm 0.53) | 75.92 (\pm 0.35) | 99.41 (\pm 0.12) | 6.7 |
| WL SUBTREE | 78.88 (\pm 0.46) | 75.45 (\pm 0.33) | 98.51 (\pm 0.05) | 4.8 |
| WL SHORTEST PATH | 75.66 (\pm 0.42) | 71.88 (\pm 0.22) | 99.36 (\pm 0.02) | 11.8 |
| WL PYRAMID MATCH | OUT-OF-MEM | 75.63 (\pm 0.49) | 99.37 (\pm 0.04) | 2.1 |
| NEIGHBORHOOD HASH | 76.02 (\pm 0.94) | 75.55 (\pm 1.00) | 99.54 (\pm 0.02) | 5.0 |
| NEIGHBORHOOD SUBGRAPH PAIRWISE DISTANCE | 78.76 (\pm 0.56) | 73.17 (\pm 0.76) | 98.04 (\pm 0.20) | 8.0 |
| ORDERED DAGs DECOMPOSITION | 75.82 (\pm 0.54) | 70.49 (\pm 0.64) | 90.75 (\pm 0.30) | 11.4 |
| PYRAMID MATCH | 76.98 (\pm 0.84) | 71.90 (\pm 0.79) | 99.56 (\pm 0.08) | 8.2 |
| GRAPHHOPPER | TIMEOUT | 74.19 (\pm 0.42) | 99.57 (\pm 0.02) | 9.6 |
| SUBGRAPH MATCHING | OUT-OF-MEM | OUT-OF-MEM | 91.96 (\pm 0.18) | 11.2 |
| PROPAGATION | 78.43 (\pm 0.55) | 72.71 (\pm 0.62) | 96.51 (\pm 0.38) | 8.4 |
| MULTISCALE LAPLACIAN | 78.28 (\pm 0.99) | 73.89 (\pm 0.93) | 98.48 (\pm 0.12) | 6.0 |
| CORE WL | 78.91 (\pm 0.50) | 75.46 (\pm 0.38) | 98.70 (\pm 0.09) | 4.1 |
| CORE SHORTEST PATH | 79.33 (\pm 0.65) | 76.31 (\pm 0.40) | 99.47 (\pm 0.05) | 5.5 |

[Nikolentzos et al., arXiv:1904.12218]

Running Time (Node-Labeled Graphs)

| KERNELS | DATASETS | | | |
|---|-----------|---------------|----------------|------------|
| | MUTAG | ENZYMEs | NCI1 | PTC-MR |
| VERTEX HISTOGRAM | 0.01s | 0.04s | 0.84s | 0.02s |
| RANDOM WALK | 1M 46.86s | 4H 24M 16.26s | TIMEOUT | 6M 41.20s |
| SHORTEST PATH | 0.92s | 11.03s | 1M 9.69s | 1.52s |
| WL SUBTREE | 0.21s | 3.81s | 7M 5.33s | 0.55s |
| WL SHORTEST PATH | 7.02s | 1M 27.07s | 15M 29.50s | 12.55s |
| WL PYRAMID MATCH | 3M 42.07s | 1H 5M 37.26s | 13H 31M 34.36s | 11M 8.16s |
| NEIGHBORHOOD HASH | 0.40s | 11.17s | 7M 4.54s | 1.31s |
| NEIGHBORHOOD SUBGRAPH PAIRWISE DISTANCE | 4.05s | 27.02s | 6M 9.81s | 7.66s |
| ORDERED DAGs DECOMPOSITION | 1.54s | 50.05s | 46M 2.13s | 4.03s |
| PYRAMID MATCH | 2.59s | 31.38s | 37M 37.50s | 11.35s |
| GRAPHHOPPER | 24.70s | 15M 38.33s | 3H 45M 8.31s | 1M 33.90s |
| SUBGRAPH MATCHING | 1M 57.25s | 3H 25M 43.59s | TIMEOUT | 4M 19.80s |
| PROPAGATION | 0.48s | 12.05s | 10M 27.83s | 1.81s |
| MULTISCALE LAPLACIAN | 10M 3.15s | 56M 43.76s | 5H 30M 56.29s | 19M 22.43s |
| CORE WL | 0.55s | 12.52s | 14M 30.56s | 17M 2.27s |
| CORE SHORTEST PATH | 2.69s | 48.02s | 3M 16.54s | 3.97s |

| KERNELS | DATASETS | | | AVG. RANK |
|---|---------------|---------------|---------------|-----------|
| | D&D | PROTEINS | AIDS | |
| VERTEX HISTOGRAM | 0.24s | 0.10s | 0.25s | 1.0 |
| RANDOM WALK | OUT-OF-MEM | 51M 10.11s | 1H 51M 56.47s | 13.6 |
| SHORTEST PATH | 55M 58.79s | 1M 18.91s | 13.93s | 4.4 |
| WL SUBTREE | 5M 52.96s | 32.48s | 40.49s | 2.8 |
| WL SHORTEST PATH | 7H 27M 21.90s | 8M 3.68s | 1M 33.46s | 10.1 |
| WL PYRAMID MATCH | OUT-OF-MEM | 5H 37M 10.33s | 5H 55M 20.37s | 14.6 |
| NEIGHBORHOOD HASH | 6M 17.21s | 41.81s | 33.30s | 3.5 |
| NEIGHBORHOOD SUBGRAPH PAIRWISE DISTANCE | 4H 36M 28.97s | 9M 9.80s | 1M 12.31s | 8.1 |
| ORDERED DAGs DECOMPOSITION | 27M 59.18s | 4M 7.81s | 2M 5.32s | 8.7 |
| PYRAMID MATCH | 5M 48.51s | 1M 26.82s | 2M 48.04s | 8.0 |
| GRAPHHOPPER | TIMEOUT | 3H 43M 1.54s | 38M 51.78s | 12.1 |
| SUBGRAPH MATCHING | OUT-OF-MEM | OUT-OF-MEM | 4H 26M 46.71s | 14.0 |
| PROPAGATION | 9M 34.30s | 51.20s | 1M 43.62s | 5.5 |
| MULTISCALE LAPLACIAN | 3H 40M 30.72s | 2H 20M 39.57s | 1H 11M 58.23s | 13.2 |
| CORE WL | 17M 2.27s | 1M 16.74s | 54.79s | 7.2 |
| CORE SHORTEST PATH | 5H 2M 39.71s | 3M 31.97s | 40.11s | 7.2 |

[Nikolentzos et al., arXiv:1904.12218]

Graph Classification (Unlabeled Graphs)

| KERNELS | DATASETS | | | | | | AVG. RANK |
|---|---------------------|---------------------|---------------------|---------------------|---------------------|---------------------|--------------|
| | IMDB BINARY | IMDB MULTI | REDDIT BINARY | REDDIT MULTI-5K | REDDIT MULTI-12K | COLLAB | |
| VERTEX HISTOGRAM | 46.54 (\pm 0.80) | 29.59 (\pm 0.40) | 47.32 (\pm 0.66) | 17.92 (\pm 0.42) | 21.73 (\pm 0.00) | 52.00 (\pm 0.00) | 12.4 |
| RANDOM WALK | 63.87 (\pm 1.06) | 45.75 (\pm 1.03) | TIMEOUT | TIMEOUT | OUT-OF-MEM | 68.00 (\pm 0.07) | 7.6 |
| SHORTEST PATH | 55.18 (\pm 1.23) | 39.37 (\pm 0.84) | 81.67 (\pm 0.23) | 47.90 (\pm 0.13) | TIMEOUT | 58.80 (\pm 0.08) | 8.3 |
| GRAPHLET | 65.19 (\pm 0.97) | 39.82 (\pm 0.89) | 76.80 (\pm 0.27) | 34.06 (\pm 0.38) | 23.08 (\pm 0.11) | 70.63 (\pm 0.25) | 7.0 |
| WL SUBTREE | 72.47 (\pm 0.50) | 50.76 (\pm 0.30) | 67.96 (\pm 1.01) | OUT-OF-MEM | OUT-OF-MEM | 78.12 (\pm 0.17) | 4.2 |
| WL SHORTEST PATH | 55.87 (\pm 1.19) | 39.63 (\pm 0.68) | TIMEOUT | TIMEOUT | TIMEOUT | 58.80 (\pm 0.06) | 10.8 |
| NEIGHBORHOOD HASH | 73.34 (\pm 0.98) | 50.68 (\pm 0.50) | 81.65 (\pm 0.28) | 49.36 (\pm 0.18) | 39.62 (\pm 0.19) | 79.99 (\pm 0.39) | 2.3 |
| NEIGHBORHOOD SUBGRAPH PAIRWISE DISTANCE | 68.81 (\pm 0.71) | 45.10 (\pm 0.63) | TIMEOUT | TIMEOUT | TIMEOUT | TIMEOUT | 7.5 |
| LOVÁSZ- ϑ | 49.21 (\pm 1.33) | 39.33 (\pm 0.95) | TIMEOUT | TIMEOUT | TIMEOUT | TIMEOUT | 15.0 |
| SVM- ϑ | 51.35 (\pm 1.54) | 38.40 (\pm 0.60) | 74.54 (\pm 0.27) | 29.65 (\pm 0.53) | 23.04 (\pm 0.18) | 55.72 (\pm 0.31) | 10.1 |
| ORDERED DAGs DECOMPOSITION | 64.70 (\pm 0.73) | 46.80 (\pm 0.51) | 50.61 (\pm 1.06) | 42.99 (\pm 0.09) | 29.83 (\pm 0.08) | 52.00 (\pm 0.00) | 7.5 |
| PYRAMID MATCH | 66.67 (\pm 1.45) | 45.25 (\pm 0.79) | 86.77 (\pm 0.42) | 48.22 (\pm 0.29) | 41.15 (\pm 0.17) | 74.57 (\pm 0.34) | 4.1 |
| GRAPHHOPPER | 57.69 (\pm 1.31) | 40.04 (\pm 0.91) | TIMEOUT | TIMEOUT | TIMEOUT | 60.21 (\pm 0.10) | 9.3 |
| SUBGRAPH MATCHING | TIMEOUT | TIMEOUT | OUT-OF-MEM | OUT-OF-MEM | OUT-OF-MEM | TIMEOUT | — |
| PROPAGATION | 51.15 (\pm 1.67) | 33.15 (\pm 1.08) | 63.41 (\pm 0.77) | 34.32 (\pm 0.61) | 24.07 (\pm 0.11) | 58.67 (\pm 0.15) | 10.1 |
| MULTISCALE LAPLACIAN | 70.94 (\pm 0.93) | 47.92 (\pm 0.87) | 89.44 (\pm 0.30) | 35.01 (\pm 0.65) | OUT-OF-MEM | 75.29 (\pm 0.49) | 3.8 |
| CORE WL | 73.31 (\pm 1.06) | 50.79 (\pm 0.54) | 72.82 (\pm 1.05) | OUT-OF-MEM | OUT-OF-MEM | OUT-OF-MEM | 3.8 |
| CORE SHORTEST PATH | 69.37 (\pm 0.68) | 50.79 (\pm 0.57) | 90.76 (\pm 0.14) | TIMEOUT | OUT-OF-MEM | TIMEOUT | 2.5 |

[Nikolentzos et al., arXiv:1904.12218]

Running Time (Unlabeled Graphs)

| KERNELS | DATASETS | | | | | | AVG. RANK |
|---|----------------|---------------|------------------|--------------------|---------------------|----------------|--------------|
| | IMDB BINARY | IMDB MULTI | REDDIT BINARY | REDDIT MULTI-5K | REDDIT MULTI-12K | COLLAB | |
| VERTEX HISTOGRAM | 0.07s | 0.15s | 0.67s | 2.20s | 6.37s | 1.12s | 1.0 |
| RANDOM WALK | 7M 20.94s | 13M 40.75s | TIMEOUT | TIMEOUT | TIMEOUT | 13H 38M 11.49s | 13.6 |
| SHORTEST PATH | 11.51s | 7.92s | 4H 48M 11.19s | 12H 40M 19.50s | TIMEOUT | 1H 9M 5.50s | 7.0 |
| GRAPHLET | 22M 45.89s | 21M 44.30s | 44M 45.42s | 44M 6.52s | 53M 14.22s | 2H 58M 1.14s | 9.5 |
| WL SUBTREE | 4.49s | 6.16s | 16M 2.65s | OUT-OF-MEM | OUT-OF-MEM | 38M 42.24s | 4.2 |
| WL SHORTEST PATH | 1M 32.66s | 1M 40.46s | TIMEOUT | TIMEOUT | TIMEOUT | 10H 27M 41.97s | 10.3 |
| NEIGHBORHOOD HASH | 21.83s | 26.07s | 23M 3.42s | 2H 44M 44.66s | 9H 11M 23.67s | 35M 49.96s | 6.3 |
| NEIGHBORHOOD SUBGRAPH PAIRWISE DISTANCE | 4M 18.12s | 2M 49.45s | TIMEOUT | TIMEOUT | TIMEOUT | TIMEOUT | 12.5 |
| LOVÁSZ- ϑ | 5H 19M 27.17s | 6H 33M 6.55s | TIMEOUT | TIMEOUT | TIMEOUT | TIMEOUT | 17.0 |
| SVM- ϑ | 39.40s | 1M 0.57s | 19M 24.73s | 23M 14.31s | 52M 10.36s | 5M 57.31s | 5.3 |
| ORDERED DAGs DECOMPOSITION | 4.47s | 4.85s | 1M 53.50s | 4M 48.92s | 8M 20.66s | 2H 1M 9.55s | 3.1 |
| PYRAMID MATCH | 1M 28.02s | 2M 13.01s | 10M 9.24s | 51M 45.10s | 3H 50M 38.60s | 36M 26.14s | 7.0 |
| GRAPHHOPPER | 2M 11.15s | 2M 3.71s | TIMEOUT | TIMEOUT | TIMEOUT | 5H 51M 32.27s | 10.3 |
| SUBGRAPH MATCHING | TIMEOUT | TIMEOUT | OUT-OF-MEM | OUT-OF-MEM | OUT-OF-MEM | TIMEOUT | — |
| PROPAGATION | 7.41s | 14.26s | 1M 23.42s | 5M 49.01s | 20M 41.73s | 4M 34.26s | 3.1 |
| MULTISCALE LAPLACIAN | 1H 22M 6.04s | 1H 41M 13.74s | 8H 21M 18.76s | 47M 51.91s | OUT-OF-MEM | 9H 24M 15.22s | 10.0 |
| CORE WL | 36.74s | 1M 1.82s | 45M 1.09s | OUT-OF-MEM | OUT-OF-MEM | OUT-OF-MEM | 8.0 |
| CORE SHORTEST PATH | 3M 58.29s | 4M 29.55s | 10H 37M 3.94s | TIMEOUT | OUT-OF-MEM | TIMEOUT | 12.3 |

[Nikolentzos et al., arXiv:1904.12218]

Graph Classification (Node-Attributed Graphs)

| KERNELS | DATASETS | | | | | AVG. RANK |
|----------------------|---------------------|---------------------|---------------------|---------------------|---------------------|--------------|
| | ENZYMES | PROTEINS_FULL | SYNTHETICNEW | SYNTHIE | BZR | |
| SHORTEST PATH | TIMEOUT | TIMEOUT | TIMEOUT | TIMEOUT | TIMEOUT | — |
| SUBGRAPH MATCHING | TIMEOUT | OUT-OF-MEM | TIMEOUT | TIMEOUT | 80.52 (\pm 0.43) | 3.0 |
| GRAPHHOPPER | 66.25 (\pm 1.24) | 72.49 (\pm 0.34) | 76.43 (\pm 1.97) | 71.75 (\pm 1.65) | 82.58 (\pm 1.05) | 1.0 |
| PROPAGATION | 15.42 (\pm 1.00) | 59.56 (\pm 0.01) | 47.90 (\pm 3.26) | 48.90 (\pm 2.05) | 78.76 (\pm 0.02) | 3.0 |
| MULTISCALE LAPLACIAN | 65.55 (\pm 0.93) | 70.55 (\pm 0.99) | 47.90 (\pm 2.13) | 69.42 (\pm 1.98) | 82.33 (\pm 1.29) | 2.0 |

[Nikolentzos et al., arXiv:1904.12218]

Running Time (Node-Attributed Graphs)

| KERNELS | DATASETS | | | | | AVG. RANK |
|----------------------|------------|---------------|---------------|------------|-------------|--------------|
| | ENZYMES | PROTEINS_FULL | SYNTHETICNEW | SYNTHIE | BZR | |
| SHORTEST PATH | TIMEOUT | TIMEOUT | TIMEOUT | TIMEOUT | TIMEOUT | — |
| SUBGRAPH MATCHING | TIMEOUT | OUT-OF-MEM | TIMEOUT | TIMEOUT | 8h 2m 3.79s | 4.0 |
| GRAPHHOPPER | 16M 36.12s | 5h 16M 46.48s | 13M 54.36s | 24M 20.00s | 4M 24.79s | 2.6 |
| PROPAGATION | 15.85s | 1M 43.58s | 13.44s | 34.68s | 10.40s | 1.0 |
| MULTISCALE LAPLACIAN | 26.05s | 4h 29M 35.69s | 2h 54M 31.22s | 15M 11.29s | 49M 33.60s | 2.4 |

[Nikolentzos et al., arXiv:1904.12218]

THANK YOU !

<http://www.lix.polytechnique.fr/dascim/>

Software and data sets:

http://www.lix.polytechnique.fr/dascim/software_datasets/

Preprint available at: <https://arxiv.org/pdf/1904.12218.pdf>